

PANRG  
Internet-Draft  
Intended status: Informational  
Expires: January 9, 2020

T. Enhardt  
TU Berlin  
C. Kraehenbuehl  
ETH Zuerich  
July 08, 2019

A Vocabulary of **Path Properties**  
draft-enhardt-panrg-path-properties-02

Mis en forme : Surlignage

Abstract

This document defines and categorizes information about Internet paths that an entity, such as a host, might have or want to have. This information is expressed as properties of paths between two hosts.

**Commentaire [Med1]:** This assumes a global connectivity. I don't think this restriction is justified (consider enterprise networks, mesh networks, etc.).

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

**Commentaire [Med2]:** Instead of reasoning about paths, what actually matters is the connectivity service(s) provided via an upstream network.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

**Commentaire [Med3]:** A host will have a partial visibility of the available path(s) (and their associated services). The full visibility cannot be established beforehand. Furthermore, such information is volatile (e.g., handover or multi-attachment with session continuity).

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

**Commentaire [Med4]:** GENERAL: the properties are dependent of the direction. Some properties will apply for inbound, outbound, or both.

This Internet-Draft will expire on January 9, 2020.

**Commentaire [Med5]:** Which excludes point-to-multipoint communications. Consider the case of an STB which want to retrieve TV channels using multicast.

Copyright Notice

Copyright (c) 2019 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

## Table of Contents

1. Introduction . . . . .	2
2. Terminology . . . . .	4
3. Domain Properties . . . . .	5
4. Backbone Properties . . . . .	6
5. Dynamic Properties . . . . .	7
6. Security Considerations . . . . .	8
7. IANA Considerations . . . . .	8
8. Informative References . . . . .	8
Acknowledgments . . . . .	9
Authors' Addresses . . . . .	9

## 1. Introduction

Because the current Internet provides an IP-based best-effort bit pipe, hosts have little information about paths to other hosts. A Path Aware Network exposes information about one or multiple paths through the network to hosts or the network infrastructure.

It is impossible to provide an exhaustive list of path properties, ~~as with every new technology and protocol, novel properties might become relevant~~. In this document, we specify a set of path properties which might be useful in the following use cases: Traffic policies, network monitoring, and path selection.

- o Traffic policies: Entities such as network operators or end users may want to define traffic policies leveraging path awareness. Such policies can allow or disallow sending traffic over specific networks or nodes, select ~~an appropriate~~ protocol depending on the capabilities of the on-path devices, or adjust protocol parameters to an existing path. An example of a traffic policy is a video streaming application choosing an (initial) video quality based on the achievable data rate, or the monetary cost of the link using a volume-based or flat-rate cost model. Another example is an enterprise network where all traffic has to go through a firewall, in which case the host needs to be aware of on-path firewalls.
- o Network monitoring: Network operators can use path properties (e.g., measured by on-path devices), to observe Quality of Service (QoS) characteristics of recent end-user traffic, and identify potential problems with their network early on, before the end-user complains.
- o Path selection: In some cases, entities can choose to use a certain path (or subset of paths) from a set of paths to achieve a specific goal. As the possible benefits of a well chosen path varies based on the goal, as a baseline, a path selection

**Commentaire [Med6]:** What is the causality effect here?

**Commentaire [Med7]:** I do still believe this is a weird term: networking is about manipulating paths, by essence!

**Commentaire [Med8]:** Difficult to parse

**Commentaire [Med9]:** This is about path selection covered in the 3rd bullet.

**Commentaire [Med10]:** I agree with this example, but I have the following comments:

The selection of the protocol is not only specific to on-path devices, but it can be the other way around: It is the selection of the protocol which may infer the devices that will be involved in a path. Consider the example of the 0-RTT Convert Protocol which is involved in a path only when invoked by a host; such invocation will lead to the use of MPTCP or TCPinc while this is not doable with the default path.

Another example of traffic policies is a connection which may be composed of multiple streams, each stream with specific requirements. A host may decide to invoke a given service function (e.g., transcoding) only for some streams while others are not processed by that service function.

**Commentaire [Med11]:** This is linked to a network, not a path.

**Commentaire [Med12]:** How this information is useful for enforcing "traffic policies" discussed in this bullet?

What would be interesting is to have means to control such firewalls.

**Commentaire [Med13]:** Which information is exposed here ?

algorithm should aim to not perform worse than the default case most of the time. Depending on the goal, an entity may prefer paths with different properties, e.g., retrieving a small webpage as quickly as possible requires low latency paths, or retrieving a large file in a peer-to-peer network requires paths with high achievable data rate. Additionally, there may be trade-offs between path properties (e.g., latency and data rate), and entities may influence these trade-offs with their choices. <sup>A</sup>

Mis en forme : Surlignage

network (e.g., an AS) can adjust its path selection for internal or external routing based on the path properties. In BGP, the Multi Exit Discriminator (MED) attribute decides which path to choose if other attributes are equal; in a path aware network, instead of using this single MED value, other properties such as maximum or available/expected data rate could additionally be used to improve load balancing. A host might be able to select between a set of paths, either if there are several paths to the same destination (e.g., if the host is a mobile device with two wireless interfaces, both providing a path), or if there are several destinations, and thus several paths, providing the same service (e.g., Application-Layer Traffic Optimization (ALTO) [RFC5693], an application layer peer-to-peer protocol allowing hosts a better-than-random peer selection). Care needs to be taken when selecting paths based on path properties, as path properties that were previously measured may have become outdated and, thus, useless to predict the path properties of packets sent now.

Such path properties may be relatively dynamic, e.g. current Round Trip Time, close to the origin, e.g. nature of the access technology on the first hop, or far from the origin, e.g. list of ASes traversed.

Usefulness over time is fundamentally different for dynamic and non-dynamic properties. The merit of a momentary measurement of a dynamic path property diminishes greatly as time goes on, e.g. the merit of an RTT measurement from a few seconds ago is quite small,

while a non-dynamic path property might stay relevant, e.g. a NAT can be assumed to stay on a path during the lifetime of a connection, as the removal of the NAT would break the connection.

Non-dynamic properties are further separated into (local) domain properties related to the first few hops of the connection, and backbone properties related to the remaining hops. Domain properties expose a high amount of information to hosts and strongly influence the connection behavior while there is little influence and information about backbone properties.

**Commentaire [Med14]:** Not for MPTCP/QUIC.

**Commentaire [Med15]:** Advanced service (e.g., NAT, IDS, ..) may be located in a Centralized PoP or hosted by a third party (DDoS protection, for example). Do you consider these as domain properties?

**Commentaire [Med16]:** Why?

**Commentaire [Med17]:** ?

**Commentaire [Med18]:** You may need to call out your topology assumptions because some conclusions may not apply to some topologies.

Dynamic properties are not separated into domain and backbone properties, since most of these properties are defined for a complete path and it is difficult and seldom useful to define them on part of the path. There are exceptions such as dynamic wireless access properties, but these do not justify separation into different categories.

This document addresses the first of the questions in Path-Aware Networking [I-D.irtf-panrg-questions], which is a product of the PANRG in the IRTF.

## 2. Terminology

**Node:** An entity which processes packets, e.g., sends, receives, forwards, or modifies them.

**Host:** ~~A node that processes packets that are explicitly addressed to itself. A node which is not router.~~

**Router:** A node that processes packets that are not explicitly addressed to itself.

**Link:** A medium or communication facility that connects two or more nodes with each other and enables them to exchange packets. A link can be physical, e.g., a Wi-Fi network which connects an Access Point to stations, or virtual, e.g., a virtual switch which connects two virtual machines hosted on the same physical machine.

**Path element:** Either a node or a link.

**Path:** A sequence of adjacent path elements, alternating between nodes and links, starting and ending with a host. A path can be viewed as an abstraction on a specific layer, omitting lower layer path elements. For example, a router implementing IPv6 may be a path element on a path when considering the network layer. If this router does not implement transport layer functionality, it is hidden when a higher layer, such as the transport or application layer, is considered. In the case of multicast or broadcast, a single packet may be sent over multiple paths at once - one path for each combination of sending and receiving host.

**Subpath:** Given a path, a subpath is a sequence of adjacent path elements of this path, starting and ending with a node.

**Flow:** One or multiple packets which are traversing the same subpath or path. For example, a flow can consist of all packets sent within a TCP session with the same five-tuple between two hosts, or it can consist of all packets sent on the same physical link.

**Commentaire [Med19]:** I personally prefer the definition in RFC1122

**Commentaire [Med20]:** This is about receiving. You may also cover sending

**Commentaire [Med21]:** This definition does cover, for example, the case of a path identified by an ordered list of AS Number.

Check for example the definition of ERO in RFC3209, or even in recent documents such as RFC7570

**Commentaire [Med22]:** This definition excludes recent notion such as SFP defined in RFC7665

**Commentaire [Med23]:** Contradicts with the sentence about alternating nodes/links.

**Property:** A trait of one or a sequence of path elements, or a trait of a flow with respect to one or a sequence of path elements. An example of a link property is the maximum data rate that can be sent over the link. An example of a node property is the administrative domain that the node belongs to. An example of a property of a flow with respect to a subpath is the aggregated one-way delay of the flow being sent from one node to another node over a subpath. A property is thus described by a tuple containing the sequence of path elements, the flow or an empty set if no packets are relevant for the property, the name of the property (e.g., maximum data rate), and the value of the property (e.g., 1Gbps).

**Commentaire [Med24]:** I'd like to double check if the defining covers also the concept of SFC/SFP. The path will be built to accommodate a property (result of SFC).

**Aggregated property:** A collection of multiple values of a property into a single value, according to a function. A property can be aggregated over multiple path elements (i.e., a path), e.g., the MTU of a path as the minimum MTU of all links on the path, over multiple packets (i.e., a flow), e.g., the median one-way latency of all packets between two nodes, or over both, e.g., the mean of the queueing delays of a flow on all nodes along a path. The aggregation function can be numerical, e.g., median, sum, minimum, it can be logical, e.g., "true if all are true", "true if at least 50% of values are true", or an arbitrary function which maps multiple input values to an output value.

**Commentaire [Med25]:** Note this cannot be applied to all properties of a connectivity service.

**Measured property:** A property that is observed for a specific path element or path, e.g., using measurements. For example, the one-way delay of a specific packet can be measured.

**Commentaire [Med26]:** I prefer « Observed » rather than measured because properties may not only be about traffic performance metrics.

**Estimated property:** An approximate calculation ~~or judgment~~ of the value of a property. For example, an estimated property may describe the expected median one-way latency of packets sent on a path within the next second. An estimated property includes the reliability of the estimate. The notion of reliability depends on the property. For example, it may be the confidence level and interval for numerical properties or the likelihood that a property holds for non-numerical properties.

**Commentaire [Med27]:** As above, this is too measurement-centric. This may not be applicable to all properties.

### 3. Domain Properties

Domain path properties relate to path elements within the first hop or the first few hops, which are usually in the same administrative domain as a host considering them.

**Commentaire [Med28]:** I'm afraid this section (and section 4) implies architectural considerations which may not be appropriate in a terminology document.

**Commentaire [Med29]:** 2, 10 ?

**Commentaire [Med30]:** L2, L3, ?

Due to the potential physical proximity and pre-existing trust or contractual relationships between hosts and path elements within the same domain, domain properties may be more accessible to the host than other properties.

**Commentaire [Med31]:** This would exclude then a host connected to a UE connected to a PLMN or a CPE connected to an enterprise or provider networks. These elements do not belong to the same administrative entity.

**Commentaire [Med32]:** Which contractual relationships would apply to this case?

**Commentaire [Med33]:** I'm not sure about this.

Furthermore, hosts may be able to influence both which domain they are in and which path elements in this domain to connect to, and they may be able to influence the properties of path elements within this domain. For example, a user might select between multiple potential adjacent links by selecting between multiple available Wi-Fi Access Points. Or when connected to an Access Point, the user may move closer to enable their device to use a different access technology, potentially increasing the data rate available to the device.

Another example is a user changing their data plan to reduce the

Monetary Cost to transmit a given amount of data across a network.

**Access Technology:** The physical or link layer technology used for transmitting or receiving a flow on one or multiple path elements in the same domain. The Access Technology may be given in an abstract way, e.g., as a Wi-Fi, Wired Ethernet, or Cellular link. It may also be given as a specific technology, e.g., as a 2G, 3G, 4G, or 5G cellular link, or an 802.11a, b, g, n, or ac WiFi link. Other path elements relevant to the access technology may include on-path devices, such as elements of a cellular backbone network. Note that there is no common registry of possible values for this property.

**Monetary Cost:** The price to be paid to transmit a specific flow across a subpath.

#### 4. Backbone Properties

Backbone path properties relate to path elements not within the same domain as a host considering them, thus, in the backbone from the host's point of view.

Typically, backbone properties are less accessible to a host than domain properties, due to the potential increased distance and the lack of pre-existing trust or contractual relationship.

Additionally, hosts are less likely to be able to influence which path elements form their path in the backbone, as well as their properties.

Some path properties relate to the entire path or to subpaths, part of which often lies outside of a host's domain. Thus, such properties are listed as Backbone Properties.

Presence of a certain ~~network-service~~ function on the path: Indicates that a

node performs a certain ~~network-service~~ function on a flow, e.g., whether

the node acts as a proxy, as a firewall, or performs Network Address Translation (NAT). This node may be either in the same domain as the host or in a different domain, i.e., the backbone.

**Commentaire [Med34]:** Connecting to a domain does not mean that the host belongs to that domain.

**Commentaire [Med35]:** An example ?

**Commentaire [Med36]:** How a host can influence the one-way delay provided by a network?

**Mis en forme :** Surlignage

**Commentaire [Med37]:** This is basically in a domain != from the local network

**Commentaire [Med38]:** A network

**Commentaire [Med39]:** Not sure about this.

**Commentaire [Med40]:** This may contradict with slicing use cases.

**Commentaire [Med41]:** To be aligned with RFC7665





Administrative Entity: The administrative entity, e.g., the AS, to which a path element or subpath belongs.

**Commentaire [Med42]:** This is an administrative domain, not \* entity \*

Disjointness: For a set of two paths, the number of shared path elements can be a measure of intersection (e.g., Jaccard coefficient, which is the number of shared elements divided by the total number of elements). Conversely, the number of non-shared path elements can be a measure of disjointness (e.g.,  $1 - \text{Jaccard coefficient}$ ). A multipath protocol might use disjointness of paths as a metric to reduce the number of single points of failure.

**Commentaire [Med43]:** you may also refer to the PCE RFCs which defines also the notion of near disjointness

Path MTU: The maximum size, in octets, of an IP packet that can be transmitted without fragmentation on a subpath.

Transport Protocols available: Whether a specific transport protocol can be used to establish a connection over a path or subpath. A host may cache its knowledge about recent successfully established connections using specific protocols, e.g., a QUIC connection, or an MPTCP subflow.

Protocol Features available: Whether a specific protocol feature is available over a path or subpath, e.g., Explicit Congestion Notification (ECN), or TCP Fast Open.

## 5. Dynamic Properties

Dynamic path properties relate to the transmission of an individual packet or of a flow over a subpath. Properties related to a path element which constitutes a single layer 2 domain are abstracted from the used physical and link layer technology, similar to [RFC8175].

Typically, Dynamic Properties can be measured or approximated, and might be made available in an aggregated form, such as averages or minimums. Dynamic Path Properties can be measured by the host itself or by a different entity. See [ANRW18-Metrics] for a discussion of how to measure some dynamic path properties at the host.

Some dynamic properties are defined in different directions for the same path element, e.g., for transmitting and receiving packets.

Maximum Data Rate (Transmit/Receive): The theoretical maximum data rate, in bits per second, that can be achieved on a link, subpath, or path, for receiving or transmitting traffic.

Current Data Rate (Transmit/Receive): The data rate, in bits per second, at which a link is currently receiving or transmitting traffic.

**Latency:** The time delay between a node sending a packet and a different node on the same path receiving the same packet.

**Latency variation:** The variation of the latency within a flow.

**Packet Loss:** The percentage of packets within a flow which are sent by one node, but not received by a different node.

**Congestion:** Whether a protocol feature such as ECN has provided information that there currently is congestion on a path.

**Commentaire [Med44]:** Please use conventional terms : one-way delay, one-way delay variation, one-way loss, etc.

## 6. Security Considerations

If nodes are basing policy or path selection decisions on path properties, they need to rely on the accuracy of path properties that other devices communicate to them. In order to be able to trust such path properties, nodes may need to establish a trust relationship or be able to verify the authenticity, integrity, and correctness of path properties received from another node.

## 7. IANA Considerations

This document has no IANA actions.

## 8. Informative References

### [ANRW18-Metrics]

Enghardt, T., Tiesel, P., and A. Feldmann, "Metrics for access network selection", Proceedings of the Applied Networking Research Workshop on - ANRW '18, DOI 10.1145/3232755.3232764, 2018.

### [I-D.irtf-panrg-questions]

Trammell, B., "Open Questions in Path Aware Networking", draft-irtf-panrg-questions-02 (work in progress), May 2019.

[RFC5693] Seedorf, J. and E. Burger, "Application-Layer Traffic Optimization (ALTO) Problem Statement", RFC 5693, DOI 10.17487/RFC5693, October 2009, <<https://www.rfc-editor.org/info/rfc5693>>.

[RFC8175] Ratliff, S., Jury, S., Satterwhite, D., Taylor, R., and B. Berry, "Dynamic Link Exchange Protocol (DLEP)", RFC 8175, DOI 10.17487/RFC8175, June 2017, <<https://www.rfc-editor.org/info/rfc8175>>.

#### Acknowledgments

Thanks to the Path-Aware Networking Research Group for the discussion and feedback. Thanks to Adrian Perrig and Matthias Rost for the feedback. Thanks to Paul Hoffman for the editorial changes.

#### Authors' Addresses

Theresa Enghardt  
TU Berlin

Email: [theresa@inet.tu-berlin.de](mailto:theresa@inet.tu-berlin.de)

Cyrill Kraehenbuehl  
ETH Zuerich

Email: [cyrill.kraehenbuehl@inf.ethz.ch](mailto:cyrill.kraehenbuehl@inf.ethz.ch)