

Network Working Group
Internet-Draft
Intended status: Standards Track
Expires: December 30, 2018

C. Filsfils, Ed.
Cisco Systems, Inc.
S. Previdi
Individual
J. Leddy
Comcast
S. Matsushima
Softbank
D. Voyer, Ed.
Bell Canada
June 28, 2018

IPv6 Segment Routing Header (SRH)
draft-ietf-6man-segment-routing-header-14

Abstract

Segment Routing can be applied to the IPv6 data plane using a new type of Routing Extension Header. This document describes the Segment Routing Extension Header and how it is used by Segment Routing capable nodes. This extension header is used within domains owned by the same administrative entity.

Requirements Language

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on December 30, 2018.

Copyright Notice

Copyright (c) 2018 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction	3
2. Segment Routing Extension Header	4
2.1. SRH TLVs	5
2.1.1. Padding TLVs	6
2.1.2. HMAC TLV	8
3. SR Nodes	9
3.1. Source SR Node	9
3.2. Transit Node	9
3.3. SR Segment Endpoint Node	9
4. Packet Processing	9
4.1. Source SR Node	9
4.1.1. Reduced SRH	10
4.2. Transit Node	10
4.3. SR Segment Endpoint Node	11
4.3.1. FIB Entry Is Locally Instantiated SRv6 END SID	11
4.3.2. FIB Entry is a Local Interface	13
4.3.3. FIB Entry Is A Non-Local Route	13
4.3.4. FIB Entry Is A No Match	13
4.3.5. Load Balancing and ECMP	13
5. Illustrations	14
5.1. Abstract Representation of an SRH	14
5.2. Example Topology	15
5.3. Source SR Node	15
5.3.1. Intra SR Domain Packet	15
5.3.2. Transit Packet Through SR Domain	16
5.4. Transit Node	16
5.5. SR Segment Endpoint Node	16
6. Security Considerations	17
6.1. Threat model	17
6.1.1. Source routing threats	17
6.1.2. Applicability of RFC 5095 to SRH	18

6.1.3.	Service stealing threat	19
6.1.4.	Topology disclosure	19
6.1.5.	ICMP Generation	19
6.2.	Security fields in SRH	19
6.2.1.	Selecting a hash algorithm	21
6.2.2.	Performance impact of HMAC	21
6.2.3.	Pre-shared key management	22
6.3.	Deployment Models	22
6.3.1.	Nodes within the SR domain	22
6.3.2.	Nodes outside of the SR domain	22
6.3.3.	SR path exposure	23
6.3.4.	Impact of BCP-38	23
7.	IANA Considerations	24
7.1.	Segment Routing Header Flags Register	24
7.2.	Segment Routing Header TLVs Register	24
8.	Implementation Status	25
8.1.	Linux	25
8.2.	Cisco Systems	25
8.3.	FD.io	25
8.4.	Barefoot	25
8.5.	Juniper	26
9.	Contributors	26
10.	Acknowledgements	26
11.	References	26
11.1.	Normative References	26
11.2.	Informative References	27
	Authors' Addresses	29

1. Introduction

Segment Routing (SR) can be applied to the IPv6 data plane using a new type of Routing Extension Header (SRH). This document describes the ~~Segment Routing Extension Header~~SRH and how it is used by Segment Routing capable nodes.

The Segment Routing ~~Architecture~~architecture [I-D.ietf-spring-segment-routing] describes Segment Routing and its instantiation in two data planes: MPLS and IPv6.

SR with the MPLS data plane is defined in [I-D.ietf-spring-segment-routing-mpls].

SR with the IPv6 data plane is defined in [I-D.filsfils-spring-srv6-network-programming].

The encoding of MPLS labels and label stacking are defined in [RFC3032].

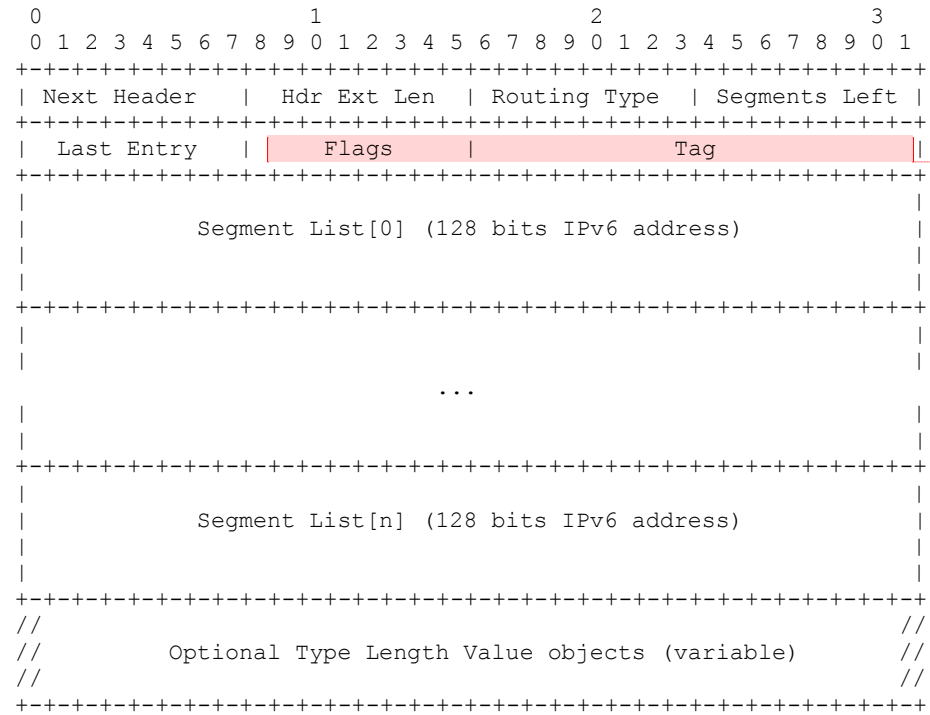
The encoding of IPv6 segments in the ~~Segment Routing Extension Header~~SRH is defined in this document.

~~Terminology Terms~~ used within this document ~~is~~are defined ~~in detail~~ in [I-D.ietf-spring-segment-routing], ~~Specifically~~specifically, these terms:
Segment Routing, SR Domain, SRv6, Segment ID (SID), SRv6 SID, Active Segment, and SR Policy.

2. Segment Routing Extension Header

Routing Headers are defined in [RFC8200]. The Segment Routing Header has a ~~new~~ Routing Type set to TBD1 (see IANA section) ~~(suggested value 4) to be assigned by IANA.~~

The format of the Segment Routing Header (SRH) is ~~defined~~ as follows:



Commentaire [Med1]: Given that no meaning is defined for each of these bits, why not define all these bits as unassigned and leave the structure of this field to future specs?

where:

- o Next Header: Defined in [RFC8200].

- o Hdr Ext Len: Defined in [RFC8200].
- o Routing Type: TBD~~1, to be assigned by IANA~~ (suggested value: 4).
- o Segments Left: Defined in [RFC8200].
- o Last Entry: contains the index (zero based), in the Segment List, of the last element of the Segment List.
- o Flags: 8 bits of flags. ~~Following flags are defined~~All the flag bits are unassigned, but a meaning can be defined in the future for these bits.

```

0 1 2 3 4 5 6 7
+---+---+---+---+
|U U U U U U U U|
+---+---+---+---+

```

U: ~~Unused~~Unassigned and for future use. MUST be 0 on transmission and ignored on receipt.

- o Tag: tag a packet as part of a class or group of packets, e.g., packets sharing the same set of properties. When 'tag' is not used at source, it MUST be set to zero on transmission. When 'tag' is not used during SRH ~~Processing~~processing it SHOULD be ignored. The allocation and use of 'tag' is outside the scope of this document.

Commentaire [Med2]: You may cite one or two examples.

Commentaire [Med3]: and/or supported.

Commentaire [Med4]: Why not MUST?

- o Segment List[n]: 128 bit IPv6 addresses representing the nth segment in the Segment List. The Segment List is encoded starting from the last segment of the SR Policy. I.e., the first element of the segment list (Segment List [0]) contains the last segment of the SR Policy, the second element contains the penultimate segment of the SR Policy and so on.

o Type Length Value optional objects (TLVs) are described in Section 2.1.

2.1. SRH TLVs

This section defines TLVs of the Segment Routing Header. The generic format of an SRH TLV is shown in Figure X.

```

0                               1
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|   Type   | Length | Variable length data |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+

```

Figure X: Generic SRH TLV format.

The description of the fields is as follows:

Type: An 8 bit value. Type values are assigned by IANA. -Unrecognized
Types MUST be ignored on receipt.

Length: The length of the ~~Variable-variable~~ length data. It is RECOMMENDED that the total length of new TLVs be multiple of 8 bytes to avoid the use of Padding TLVs.

Variable length data: Length bytes of data that is specific to the Type.

~~Type Length Value (TLV)~~A TLV contains ~~OPTIONAL-optional~~ information that may be used by the node identified in the Destination Address (DA) of the packet.

Each TLV has its own length, format and semantic. ~~The code point allocated (by IANA) to each TLV Type defines both the format and the semantic of the information carried in the TLV.~~ Multiple TLVs may be encoded in the same SRH. Whether the same TLV is present multiple times within the same SRH is specified in the corresponding TLV specification. Likewise, order constraints, if any, are specified in those documents.

Commentaire [Med5]: Redundant with the text above.

TLVs may change en route at each segment. To identify when a TLV type may change en route the most significant bit of the Type has the following significance:

Mis en forme : Surlignage

0: TLV data does not change en route

1: TLV data does change en route

Identifying which TLVs change en route, without having to understand the Type, is required for Authentication Header Integrity Check Value (ICV) computation. Any TLV that changes en route is considered mutable for the purpose of ICV computation, the Type Length and Variable Length Data is ignored for the purpose of ICV Computation as defined in [RFC4302].

Mis en forme : Surlignage

The "Length" field of the TLV is used to skip the TLV while inspecting the SRH in case the node doesn't support or recognize the Type. The "Length" defines the TLV length in octets, not including the "Type" and "Length" fields.

The following TLVs are defined in this document:

Padding TLVs

HMAC TLV

Additional TLVs may be defined in the future.

Commentaire [Med6]: Are those mandatory to implement/support?

I guess the answer is not the same for both TLVs.

The text should then be explicit about this.

The set of TLVs that are needed to be inserted and processed within a domain is deployment-specific. Such information is communicated using means that are out of the scope of the document.

2.1.1. Padding TLVs

There are two types of padding TLVs, pad0 and padN, the following applies to both:

Padding TLVs are used to pad the TLVs to a multiple of 8 octets.

More than one Padding TLV MUST NOT appear in the SRH.

~~The Padding TLVs are used to align the SRH total length on the 8 octet boundary.~~

When present, a single Pad0 or PadN TLV MUST appear as the last TLV.

When present, a PadN TLV MUST have a length from 0 to 5 in order to align the SRH total length on a 8-octet boundary.

Padding TLVs ~~are~~ MUST be ignored by a node processing the SRH TLV, even if more than one is present.

~~Padding TLVs are ignored during ICV calculation.~~

Commentaire [Med7]: This is redundant with the meaning the most significant bit.

2.1.1.1. PAD0

```

0 1 2 3 4 5 6 7
+-+--+--+--+--+
|   Type   |
+-+--+--+--+--+

```

Commentaire [Med8]: This is not a TLV.

Type: ~~to be assigned by IANA~~ TBD2 (Suggested value 128)

A single Pad0 TLV MUST be used when a single byte of padding is required. If more than one byte of padding is required a Pad0 TLV MUST NOT be used, the PadN TLV MUST be used instead.

2.1.1.2. PADN

```

      0              1              2              3
0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|   Type   | Length | Padding (variable) |
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
//                Padding (variable)                //
+-+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+

```

Type: TBD3~~to be assigned by IANA~~ (suggested value 129).


Length: May be set to a value ranging from 0 to 5.

Padding: Length octets of padding. Padding bits have no semantic. They MUST be set to 0 on transmission and ignored on receipt.

The PadN TLV MUST be used when more than one byte of padding is required.

2.1.2. HMAC TLV

HMAC TLV is OPTIONAL and contains the HMAC information. The HMAC TLV has the following format:



where:

- o Type: ~~to be assigned by IANA~~TBD4 (suggested value 5).
- o Length: 38.
- o ~~RESERVED~~Unassigned: 2 octets. MUST be set to 0 on transmission and MUST be ignored on receipt.
- o HMAC Key ID: 4 octets.
- o HMAC: 32 octets.

~~o~~ HMAC and HMAC Key ID usage is described in Section 6

~~The Following applies to the HMAC TLV:~~

~~o~~ Local policy within a domain determines when to check for an HMAC and potentially a requirement on where the HMAC TLV must appear (e.g., first TLV). This How such local policy is defined and communicated to involved nodes is are outside the scope of this document. It may be based on the active segment at an SR Segment endpoint node, the result of an ACL that considers incoming interface, or other packet fields.

Commentaire [Med9]: More text is required to describe those fields.

3. SR Nodes

There are different types of nodes that may be involved in ~~segment routing-SR~~ networks: source SR nodes originate packets with a segment in

the destination address of the IPv6 header, transit nodes that forward packets destined to a remote segment, and SR segment endpoint nodes that process a local segment in the destination address of an IPv6 header.

3.1. Source SR Node

A ~~Source source~~ SR ~~Node node~~ is any node that originates an IPv6 packet with a segment (i.e., SRv6 SID) in the destination address of the IPv6 header. The source SR node may not be the original source of the packet. — The packet leaving the source SR Node node may or may not contain an SRH. This includes either:

A host originating an IPv6 packet.

An SR domain ingress router encapsulating a received packet in an outer IPv6 header, followed by an ~~optional~~-SRH. Classification operations to decide which packets need to be stamped with an SRH are deployment-specific and thus out of scope.

The mechanism through which a segment in the destination address of the IPv6 header and the Segment List in the SRH₇ is derived₇ is outside the scope of this document.

3.2. Transit Node

A transit node is any node forwarding an IPv6 packet where the destination address of that packet is not locally configured as a segment nor a local interface. A transit node is not required to be capable of processing a segment nor SRH.

Routers MUST NOT process the SRH unless they are explicitly configured to do so. Processing the SRH MUST be disabled by default. Typically, routers that are not within an SRH-enabled domain MUST NOT be configured to process the SRH.

3.3. SR Segment Endpoint Node

A SR segment endpoint node is any node receiving an IPv6 packet where the destination address of that packet is locally configured as a segment or local interface.

4. Packet Processing

This section describes SRv6 packet processing at the SR source, Transit₇ and SR segment endpoint nodes.

Commentaire [Med10]: Why then calling it SR source node?

4.1. Source SR Node

| A ~~Source~~source SR node steers a packet into an SR Policy. If the SR Policy results in a segment list containing a single segment, and there is

no need to add information to an SRH flag or TLV, the DA is set to the single segment list entry and the SRH MAY be omitted.

When needed, the SRH is created as follows:

Next Header and Hdr Ext Len fields are set as specified in [RFC8200].

Routing Type field is set as TBD ~~(to be allocated by IANA, suggested value 4)~~.

The DA of the packet is set with the value of the first segment.

The first element of the SRH Segment List is the ultimate segment. The second element is the penultimate segment, and so on.

The Segments Left field is set to n-1 where n is the number of elements in the SR Policy.

The Last Entry field is set to n-1 where n is the number of elements in the SR Policy.

HMAC TLV may be set according to Section 6.

The packet is forwarded toward the packet's Destination Address (the first segment).

4.1.1. Reduced SRH

When a source does not require the entire SID list to be preserved in the SRH, a reduced SRH may be used.

A reduced SRH does not contain the first segment of the related SR Policy (the first segment is the one already in the DA of the IPv6 header), and the Last Entry field is set to n-2 where n is the number of elements in the SR Policy.

4.2. Transit Node

As specified in [RFC8200], the only node allowed to inspect the Routing Extension Header (and therefore the SRH), is the node corresponding to the DA of the packet. Any other transit node MUST NOT inspect the underneath routing header and MUST forward the packet toward the DA according to its IPv6 routing table.

When a SID is in the destination address of an IPv6 header of a packet, it's routed through an IPv6 network as an IPv6 address. SIDs, or the prefix(es) covering SIDs, and their reachability may be

Commentaire [Med11]: What does that mean? Do mean that is used as IPv6 address lookup?

distributed by means outside the scope of this document. For example, [RFC5308] or [RFC5340] may be used to advertise a prefix covering the SIDs on a node.

4.3. SR Segment Endpoint Node

Without constraining the details of an implementation, the SR segment endpoint node creates Forwarding Information Base (FIB) entries for its local SIDs.

When an SRv6-capable node receives an IPv6 packet, it performs a longest-prefix-match lookup on the packets destination address. This lookup can return any of the following:

- A FIB entry that represents a locally instantiated SRv6 SID
- A FIB entry that represents a local interface, not locally instantiated as an SRv6 SID
- A FIB entry that represents a non-local route
- No Match

4.3.1. FIB Entry Is Locally Instantiated SRv6 END SID

This document, and section, defines a single SRv6 SID called END. Future documents may define additional SRv6 SIDs. In which case, the entire content of this section will be defined in that document.

If the FIB entry represents a locally instantiated SRv6 SID, process the next header of the IPv6 header as defined in ~~section~~section 4 of [RFC8200]

The following sections describe the actions to take while processing next header fields.

4.3.1.1. SRH Processing

Commentaire [Med12]: Do you mean updated? or that a content following the the structure of this section will be required?


```

When an SRH is processed {
  If Segments Left is equal to zero {
    Proceed to process the next header in the packet, whose type
    is identified by the Next Header field in the Routing header.
  }
  Else {
    If local policy requires TLV processing {
      Perform TLV processing (see TLV Processing)
    }
    max_last_entry = ( Hdr Ext Len / 2 ) - 1

    If ((Last Entry > max_last_entry) or
        (Segments Left is greater than (Last Entry+1)) {
      Send an ICMP Parameter Problem, Code 0, message to the
      Source Address, pointing to the Segments Left field, and
      discard the packet.
    }
    Else {
      Decrement Segments Left by 1.
      Copy Segment List[Segments Left] from the SRH to the
      destination address of the IPv6 header.
      If the IPv6 Hop Limit is less than or equal to 1 {
        Send an ICMP Time Exceeded -- Hop Limit Exceeded in
        Transit message to the Source Address and discard
        the packet.
      }
      Else {
        Decrement the Hop Limit by 1
        Resubmit the packet to the IPv6 module for transmission
        to the new destination.
      }
    }
  }
}

```

Commentaire [Med13]: This assumes that TLV is supported.

Commentaire [Med14]: but that section does not specify the behavior.

This is IMHO specific to each TLV specification.

4.3.1.1.1. TLV Processing

Local policy determines how ~~TLV's~~ TLVs are to be processed when the Active

Segment is a local END SID. The definition of local policy and how this is communicated are ~~is~~ outside the scope of this document.

For illustration purposes only, two ~~example~~ sample local policies that may be associated with an END SID are provided below.

Example 1:

For any packet received from interface I2
Skip TLV processing

Example 2:

For any packet received from interface I1
If first TLV is HMAC {
 Process the HMAC TLV
}
Else {
 Discard the packet
}

4.3.1.2. Upper-layer Header or No Next Header

Send an ICMP destination unreachable to
the Source Address and discard the packet.

4.3.2. FIB Entry is a Local Interface

If the FIB entry represents a local interface, not locally
instantiated as an SRv6 SID, the SRH is processed as follows:

If Segments Left is zero, the node must ignore the Routing header
and proceed to process the next header in the packet, whose type
is identified by the Next Header field in the Routing Header.

If Segments Left is non-zero, the node must discard the packet and
send an ICMP Parameter Problem, Code 0, message to the packet's
Source Address, pointing to the unrecognized Routing Type.

4.3.3. FIB Entry Is A Non-Local Route

Processing is not changed by this document.

4.3.4. FIB Entry Is A No Match

Processing is not changed by this document.

4.3.5. Load Balancing and ECMP

Within an SR domain, an SR source node encapsulates a packet ~~in using~~
an outer IPv6 header for transport to an endpoint (e.g., egress node).

The SR source node

MUST impose a flow label computed based on the inner packet. The
computation of the flow label is as recommended in [RFC6438] for the
sending Tunnel End Point.

Commentaire [Med15]: this may not
be the ultimate destination node.
You may explicit this in the text.

Commentaire [Med16]: Do you mean
the SR source node?

At any transit node within an SR domain, the flow label MUST be used as defined in [RFC6438] to calculate the ECMP hash toward the destination address. If flow label is not used, the transit node may hash all packets between a pair of SR ~~Edge-edge~~ nodes to the same link.

At an SR segment endpoint node, the flow label MUST be used as defined in [RFC6438] to calculate any ECMP hash used to forward the processed packet to the next segment.

5. Illustrations

This section provides illustrations of SRv6 packet processing at SR source, transit, and SR segment endpoint nodes.

5.1. Abstract Representation of an SRH

For a node k , its IPv6 address is represented as A_k , its SRv6 SID is represented as S_k .

IPv6 headers are represented as the tuple of (source, destination). For example, a packet with source address A_1 and destination address A_2 is represented as (A_1, A_2) . The payload of the packet is omitted.

An SR Policy is a list of segments. A list of segments is represented as $\langle S_1, S_2, S_3 \rangle$ where S_1 is the first SID to visit, S_2 is the second SID to visit and S_3 is the last SID to visit.

$(SA, DA) (S_3, S_2, S_1; SL)$ represents an IPv6 packet with:

- o Source Address is SA , Destination Addresses is DA , and next-header is SRH.
- o SRH with SID list $\langle S_1, S_2, S_3 \rangle$ with $SegmentsLeft = SL$.
- o Note the difference between the $\langle \rangle$ and $()$ symbols. $\langle S_1, S_2, S_3 \rangle$ represents a SID list where the leftmost segment is the first segment. Whereas, $(S_3, S_2, S_1; SL)$ represents the same SID list but encoded in the SRH Segment List format where the leftmost segment is the last segment. When referring to an SR policy in a high-level use-case, it is simpler to use the $\langle S_1, S_2, S_3 \rangle$ notation. When referring to an illustration of detailed behavior, the $(S_3, S_2, S_1; SL)$ notation is more convenient.

At its SR Policy headend, the Segment List $\langle S_1, S_2, S_3 \rangle$ results in SRH $(S_3, S_2, S_1; SL=2)$ represented fully as:

```

Segments Left=2
Last Entry=2
Flags=0
Tag=0
Segment List[0]=S3
Segment List[1]=S2
Segment List[2]=S1

```

5.2. Example Topology

The ~~following~~ topology depicted in Figure X is used in the examples below~~.~~.




Figure X: xxxx

- o 3 and 4 are SR Domain edge routers.
- o 5, 6, and 7 are all SR Domain routers.
- o 8 and 9 are hosts within the SR Domain.
- o 1 and 2 are hosts outside the SR Domain.

5.3. Source SR Node

5.3.1. Intra SR Domain Packet

When host 8 sends a packet to host 9 via an SR Policy <S7,A9> the packet ~~includes the following information:~~

P1: (A8,S7) (A9,S7; SL=1)

5.3.1.1. Reduced Variant

When host 8 sends a packet to host 9 via an SR Policy <S7,A9> and it wants to use a reduced SRH, the SRH information enclosed in the packet is:

P2: (A8,S7) (A9; SL=1)

5.3.2. Transit Packet Through SR Domain

When host 1 sends a packet to host 2, the packet is

P3: (A1,A2)

The SR Domain ingress router 3 receives P3 and steers it to SR Domain egress router 4 via an SR Policy <S7, S4>. Router 3 encapsulates the received packet P3 in an outer header with an SRH. The packet is

P4: (A3,~~S7~~) (S4, S7; SL=1) (A1, A2)

If the SR Policy contains only one segment (the egress router 4), the ingress Router 3 encapsulates P3 into an outer header (A3, S4). The packet is

P5: (A3,~~S4~~) (A1, A2)

5.3.2.1. Reduced Variant

The SR Domain ingress router 3 receives P3 and steers it to SR Domain egress router 4 via an SR Policy <S7, S4>. If router 3 wants to use a reduced SRH, Router 3 encapsulates the received packet P3 in an outer header with a reduced SRH. The packet is

P6: (A3, S7) (S4; SL=1) (A1, A2)

5.4. Transit Node

Nodes 5 acts as transit nodes for packet P1, and sends packet

P1: (A8,S7) (A9,S7;SL=1)

on the interface toward node 7.

5.5. SR Segment Endpoint Node

Node 7 receives packet P1 and, using the logic in section 4.3.1, sends packet

P7: (A8,A9) (A9,S7; SL=0)

on the interface toward router 6.

6. Security Considerations

This section analyzes the security threat model, the security issues and proposed solutions related to the ~~new~~ Segment Routing Header.

The ~~Segment Routing Header (SRH)~~ is ~~simply~~ another type of the Routing Header as described in [RFC8200] and is:

- o Added by an SR edge router when entering the segment routing domain or by the originating host itself. The source host can ~~even~~ be outside the SR domain;
- o inspected and acted upon when reaching the destination address of the IP header per [RFC8200].

Per [RFC8200], routers on the path that ~~simply~~ forward an IPv6 packet (i.e., the IPv6 destination address is none of theirs) will never inspect and process the content of the SRH. Routers whose FIB contains a locally instantiated SRv6 SID equal to the destination address field of the IPv6 packet MUST parse the SRH if present, and if supported and if the local configuration allows it, MUST act accordingly to the SRH content.

As specified in [RFC8200], the default behavior of a non SR-capable router upon receipt of an IPv6 packet with SRH destined to an address of its, is to:

- o ignore the SRH completely if the Segment Left field is 0 and proceed to process the next header in the IPv6 packet;
- o discard the IPv6 packet if Segment Left field is greater than 0, it MAY send a Parameter Problem ICMP message back to the Source Address.

6.1. Threat model

6.1.1. Source routing threats

Using an SRH is similar to source routing, therefore it has some well-known security issues as described in [RFC4942] section 2.1.1 and [RFC5095]:

- o amplification attacks: where a packet could be forged in such a way to cause looping among a set of SR-enabled routers causing unnecessary traffic, hence a Denial of Service (DoS) against bandwidth;

Commentaire [Med17]: I suggest to add this sentence:

Security considerations discussed in [RFC7045] and [RFC7112] apply.

Commentaire [Med18]: This is part of the SRH processing describe above. No need to repeat it here.

- o reflection attack: where a hacker could force an intermediate node to appear as the immediate attacker, hence hiding the real attacker from naive forensic;
- o bypass attack: where an intermediate node could be used as a stepping stone (for example in a De-Militarized Zone) to attack another host (for example in the datacenter or any back-end server).

6.1.2. Applicability of RFC 5095 to SRH

First of all, the reader must remember this specific part of section 1 of [RFC5095], "A side effect is that this also eliminates benign RH0 use-cases; however, such applications may be facilitated by future Routing Header specifications.". In short, it is not forbidden to create new secure type of Routing Header; for example, [RFC6554] (RPL) also creates a new Routing Header type for a specific application confined in a single network.

In the segment routing architecture described in [I-D.ietf-spring-segment-routing] there are basically two kinds of nodes (routers and hosts):

- o nodes within the SR domain, which is within one single administrative domain, i.e., where all nodes are trusted anyway else the damage caused by those nodes could be worse than amplification attacks: traffic interception, man-in-the-middle attacks, more server DoS by dropping packets, and so on.
- o nodes outside of the SR domain, which is outside of the administrative segment routing domain hence they cannot be trusted because there is no physical security for those nodes, i.e., they can be replaced by hostile nodes or can be coerced in wrong behaviors.

The main use case for SR consists of the single administrative domain where only trusted nodes with SR enabled and configured participate in SR: this is the same model as in [RFC6554]. All non-trusted nodes do not participate as either SR processing is not enabled by default or because they only process SRH from nodes within their domain.

Moreover, all SR nodes ignore SRH created by outsiders based on topology information (received on a peering or internal interface) or on presence and validity of the HMAC field. Therefore, if intermediate nodes ONLY act on valid and authorized SRH (such as within a single administrative domain), then there is no security threat similar to RH-0. Hence, the [RFC5095] attacks are not applicable.

6.1.3. Service stealing threat

Segment routing is used for added value services, there is also a need to prevent non-participating nodes to use those services; this is called 'service stealing prevention'.

6.1.4. Topology disclosure

The SRH may also contains SIDs of some intermediate SR-nodes in the path towards the destination, this obviously reveals those addresses to the potentially hostile attackers if those attackers are able to intercept packets containing SRH. On the other hand, if the attacker can do a traceroute whose probes will be forwarded along the SR Policy, then there is little learned by intercepting the SRH itself.

6.1.5. ICMP Generation

Per Section 4.4 of [RFC8200], when destination nodes (i.e. where the destination address is one of theirs) receive a Routing Header with unsupported Routing Type, the required behavior is:

- o If Segments Left is zero, the node must ignore the Routing Header and proceed to process the next header in the packet.
- o If Segments Left is non-zero, the node must discard the packet and send an ICMP Parameter Problem, Code 0, message to the packet's Source Address, pointing to the unrecognized Routing Type.

This required behavior could be used by an attacker to force the generation of ICMP message by any node. The attacker could send packets with SRH (with Segment Left not set to 0) destined to a node not supporting SRH. Per [RFC8200], the destination node could generate an ICMP message, causing a local CPU utilization and if the source of the offending packet with SRH was spoofed could lead to a reflection attack without any amplification.

It must be noted that this is a required behavior for any unsupported Routing Type and not limited to SRH packets. So, it is not specific to SRH and the usual rate limiting for ICMP generation is required anyway for any IPv6 implementation and has been implemented and deployed for many years.

6.2. Security fields in SRH

This section summarizes the use of specific fields in the SRH. They are based on a key-hashed message authentication code (HMAC).

The security-related fields in the SRH are instantiated by the HMAC TLV, containing:

- o HMAC Key-id, 32 bits wide;
- o HMAC, 256 bits wide (optional, exists only if HMAC Key-id is not 0).

The HMAC field is the output of the HMAC computation (per [RFC2104]) using a pre-shared key identified by HMAC Key-id and of the text which consists of the concatenation of:

- o the source IPv6 address;
- o Last Entry field;
- o an octet of bit flags;
- o HMAC Key-id;
- o all addresses in the Segment List.

The purpose of the HMAC TLV is to verify the validity, the integrity and the authorization of the SRH itself. If an outsider of the SR domain does not have access to a current pre-shared secret, then it cannot compute the right HMAC field and the first SR router on the path processing the SRH and configured to check the validity of the HMAC will simply reject the packet.

The HMAC TLV is located at the end of the SRH simply because only the router on the ingress of the SR domain needs to process it, then all other SR nodes can ignore it (based on local policy) because they trust the upstream router. This is to speed up forwarding operations because SR routers which do not validate the SRH do not need to parse the SRH until the end.

The HMAC Key-id field allows for the simultaneous existence of several hash algorithms (SHA-256, SHA3-256 ... or future ones) as well as pre-shared keys. The HMAC Key-id field is opaque, i.e., it has neither syntax nor semantic except as an index to the right combination of pre-shared key and hash algorithm and except that a value of 0 means that there is no HMAC field. Having an HMAC Key-id field allows for pre-shared key roll-over when two pre-shared keys are supported for a while when all SR nodes converged to a fresher pre-shared key. It could also allow for interoperation among different SR domains if allowed by local policy and assuming a collision-free HMAC Key Id allocation.

When a specific SRH is linked to a time-related service (such as turbo-QoS for a 1-hour period) where the DA, Segment ID (SID) are identical, then it is important to refresh the shared-secret frequently as the HMAC validity period expires only when the HMAC Key-id and its associated shared-secret expires.

6.2.1. Selecting a hash algorithm

The HMAC field in the HMAC TLV is 256 bit wide. Therefore, the HMAC MUST be based on a hash function whose output is at least 256 bits. If the output of the hash function is 256, then this output is simply inserted in the HMAC field. If the output of the hash function is larger than 256 bits, then the output value is truncated to 256 by taking the least-significant 256 bits and inserting them in the HMAC field.

SRH implementations can support multiple hash functions but MUST implement SHA-2 [FIPS180-4] in its SHA-256 variant.

NOTE: SHA-1 is currently used by some early implementations used for quick interoperations testing, the 160-bit hash value must then be right-hand padded with 96 bits set to 0. The authors understand that this is not secure but is ok for limited tests.

6.2.2. Performance impact of HMAC

While adding an HMAC to each and every SR packet increases the security, it has a performance impact. Nevertheless, it must be noted that:

- o the HMAC field is used only when SRH is added by a device (such as a home set-top box) which is outside of the segment routing domain. If the SRH is added by a router in the trusted segment routing domain, then, there is no need for an HMAC field, hence no performance impact.
- o when present, the HMAC field need only be checked and validated by the first router of the segment routing domain, this router is named 'validating SR router'.
- o this validating SR router can also have a cache of <IPv6 header + SRH, HMAC field value> to improve the performance. It is not the same use case as in IPsec where HMAC value was unique per packet, in SRH, the HMAC value is unique per flow.
- o hash functions such as SHA-2 have been optimized for security and performance and there are multiple implementations with good performance.

With the above points in mind, the performance impact of using HMAC is minimized.

6.2.3. Pre-shared key management

The field HMAC Key-id allows for:

- o key roll-over: when there is a need to change the key (the hash pre-shared secret), then multiple pre-shared keys can be used simultaneously. The validating SR router can have a table of <HMAC Key-id, pre-shared secret> for the currently active and future keys.
- o different algorithms: by extending the previous table to <HMAC Key-id, hash function, pre-shared secret>, the validating SR router can also support several hash algorithms (see section Section 6.2.1)

The pre-shared secret distribution can be done:

- o in the configuration of the validating SR routers, either by static configuration or any SDN oriented approach;
- o dynamically using a trusted key distribution such as [RFC6407]

The intent of this document is NOT to define yet-another-key-distribution-protocol.

6.3. Deployment Models

6.3.1. Nodes within the SR domain

SR Source Nodes within an SR Domain are trusted to generate IPv6 packets with SRH. SR segment endpoint nodes receiving packets on interface that are part of the SR Domain may process any packet destined to a local segment, containing an SRH.

6.3.2. Nodes outside of the SR domain

Nodes outside the SR Domain cannot be trusted. SR Domain Ingress routers SHOULD discard packets destined to SIDs within the SR Domain (regardless of the presence of an SRH) to avoid attacks on the SR Domain. This is accomplished via infrastructure Access Lists (iACLs) applied on domain ingress nodes. However the SR Domain may be extended to nodes outside of it via use of the SRH HMAC.

Nodes outside the SR Domain may request, by some trusted means outside the scope of this document, a complete SRH including an HMAC TLV which is computed correctly for the SRH (see Section 6.2).

SR Domain ingress routers permit traffic destined to select SIDs with local policy requiring HMAC TLV processing for those select SIDs, i.e. these SIDs provide a gateway to the SR Domain for a set of segment lists.

If HMAC validation is successful, the packet is forwarded to the next segment. Within the SR Domain no further HMAC check need be performed. However, other segments in the SR domain MAY verify the HMAC TLV when the SRH is processed, dependent on local policy.

If HMAC validation fails an ICMP error message (parameter problem) SHOULD be generated (but rate limited) and SHOULD be logged.

6.3.3. SR path exposure

The SRH contains a Segment List. If an observer outside the SR Domain is able to inspect the SRH, they may use the segments in the Segment List to launch an attack on the SR Domain or obtain knowledge of the topology within the SR Domain. When the SR Source node is outside the SR Domain and the packet traverses the public internet to the SR Domain ingress router it is likely that others will have access to the Segment List in the SRH.

IPSec Encapsulating Security Payload (ESP), [RFC4303], cannot be used to protect the SRH as the ESP header must appear after the routing header (including SRH).

Exposure of segments and TLV content to observers outside the SR Domain should be considered in any deployment. There are two methods to limit exposure, and attacks on segments within the SR Domain from outside the SR Domain:

- Limit the number of segments and the TLV data exposed in SRH from nodes outside the SR Domain.

- Restrict which SIDs may accept traffic from outside the SR Domain to only those enforcing HMAC verification by using iACLs (as described in Section 6.3.2).

Mis en forme : Anglais (États Unis)

6.3.4. Impact of BCP-38

BCP-38 [RFC2827], also known as "Network Ingress Filtering", checks whether the source address of packets received on an interface is valid for this interface. The use of loose source routing such as

SRH forces packets to follow a path which differs from the expected routing. Therefore, if BCP-38 was implemented in all routers inside the SR domain, SR packets could be received by an interface which is not the expected one, and the packets could be dropped.

As BCP-38 is only deployed at the ingress routers of an administrative domain, and as Packets arriving at those ingress routers have been forwarded using the routing information, then there is no reason why this ingress router should drop the SRH packet based on BCP-38. Routers inside the domain commonly do not apply BCP-38; so, this is not a problem.

7. IANA Considerations

This document makes the following registrations in the Internet Protocol Version 6 (IPv6) Parameters "Routing Type" registry maintained by IANA:

Suggested Value	Description	Reference
TBD1 (4)	Segment Routing Header (SRH)	This document

Mis en forme : Anglais (États Unis)

Mis en forme : Anglais (États Unis)

This document request IANA to create and maintain a new Registry: "Segment Routing Header TLVs"

7.1. Segment Routing Header Flags Register

This document requests the creation of a new IANA managed registry to identify SRH Flags ~~Bits~~bits. The registration procedure is "Expert Review" as defined in [RFC8126]. Suggested registry name is "Segment Routing Header Flags". Flags ~~filed~~is 8 bits, the following bits are defined in this document:

Mis en forme : Anglais (États Unis)

Mis en forme : Anglais (États Unis)

Suggested Bit	Description	Reference
---------------	-------------	-----------

4	HMAC	This document
---	------	---------------

Commentaire [Med19]: Do you mean the position of the bit?

Mis en forme : Anglais (États Unis)

Commentaire [Med20]: This is not discussed in the core text.

Mis en forme : Anglais (États Unis)

7.2. Segment Routing Header TLVs Register

This document requests the creation of a new IANA managed registry to identify SRH TLVs. The registration procedure is "Expert Review" as defined in [RFC8126]. Suggested registry name is "Segment Routing Header TLVs". A TLV is identified through an unsigned 8 bit codepoint value. The following codepoints are defined in this document:

Suggested Value	Description	Reference

TBD4 (5)	HMAC TLV	This document
TBD2 (128)	Pad0 TLV	This document
TBD3 (129)	PadN TLV	This document

8. Implementation Status

This section is to be removed prior to publishing as an RFC.

8.1. Linux

Name: Linux Kernel v4.14

Status: Production

Implementation: adds SRH, performs END processing, supports HMAC TLV

Details: <https://irtf.org/anrw/2017/anrw17-final3.pdf> and [I-D.filsfils-spring-srv6-interop]

8.2. Cisco Systems

Name: IOS XR and IOS XE

Status: Pre-production

Implementation: adds SRH, performs END processing, no TLV processing

Details: [I-D.filsfils-spring-srv6-interop]

8.3. FD.io

Name: VPP/Segment Routing for IPv6

Status: Production

Implementation: adds SRH, performs END processing, no TLV processing

Details: https://wiki.fd.io/view/VPP/Segment_Routing_for_IPv6 and [I-D.filsfils-spring-srv6-interop]

8.4. Barefoot

Name: Barefoot Networks Tofino NPU

Status: Prototype

Implementation: performs END processing, no TLV processing

Details: [I-D.filsfils-spring-srv6-interop]

8.5. Juniper

Name: Juniper Networks Trio and vTrio NPU's

Status: Prototype & Experimental

Implementation: SRH insertion mode, Process SID where SID is an interface address, no TLV processing

9. Contributors

Kamran Raza, Darren Dukes, Brian Field, Daniel Bernier, Ida Leung, Jen Linkova, Ebben Aries, Tomoya Kosugi, Eric Vyncke, David Lebrun, Dirk Steinberg, Robert Raszuk, Dave Barach, John Brzozowski, Pierre Francois, Nagendra Kumar, Mark Townsley, Christian Martin, Roberta Maglione, James Connolly, Aloys Augustin contributed to the content of this document.

10. Acknowledgements

The authors would like to thank Ole Troan, Bob Hinden, Ron Bonica, Fred Baker, Brian Carpenter, Alexandru Petrescu, Punit Kumar Jaiswal, and David Lebrun for their comments to this document.

11. References

11.1. Normative References

[FIPS180-4]

National Institute of Standards and Technology, "FIPS 180-4 Secure Hash Standard (SHS)", March 2012, <<http://csrc.nist.gov/publications/fips/fips180-4/fips-180-4.pdf>>.

[I-D.ietf-spring-segment-routing]

Filsfils, C., Previdi, S., Ginsberg, L., Decraene, B., Litkowski, S., and R. Shakir, "Segment Routing Architecture", draft-ietf-spring-segment-routing-15 (work in progress), January 2018.

[RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.

- [RFC4303] Kent, S., "IP Encapsulating Security Payload (ESP)", RFC 4303, DOI 10.17487/RFC4303, December 2005, <<https://www.rfc-editor.org/info/rfc4303>>.
- [RFC5095] Abley, J., Savola, P., and G. Neville-Neil, "Deprecation of Type 0 Routing Headers in IPv6", RFC 5095, DOI 10.17487/RFC5095, December 2007, <<https://www.rfc-editor.org/info/rfc5095>>.
- [RFC6407] Weis, B., Rowles, S., and T. Hardjono, "The Group Domain of Interpretation", RFC 6407, DOI 10.17487/RFC6407, October 2011, <<https://www.rfc-editor.org/info/rfc6407>>.
- [RFC8174] Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <<https://www.rfc-editor.org/info/rfc8174>>.
- [RFC8200] Deering, S. and R. Hinden, "Internet Protocol, Version 6 (IPv6) Specification", STD 86, RFC 8200, DOI 10.17487/RFC8200, July 2017, <<https://www.rfc-editor.org/info/rfc8200>>.

11.2. Informative References

- [I-D.filsfils-spring-srv6-interop]
Filsfils, C., Clad, F., Camarillo, P., Abdelsalam, A., Salsano, S., Bonaventure, O., Horn, J., and J. Liste, "SRv6 interoperability report", draft-filsfils-spring-srv6-interop-00 (work in progress), March 2018.
- [I-D.filsfils-spring-srv6-network-programming]
Filsfils, C., Li, Z., Leddy, J., daniel.voyer@bell.ca, d., daniel.bernier@bell.ca, d., Steinberg, D., Raszuk, R., Matsushima, S., Lebrun, D., Decraene, B., Peirens, B., Salsano, S., Naik, G., Elmalky, H., Jonnalagadda, P., and M. Sharif, "SRv6 Network Programming", draft-filsfils-spring-srv6-network-programming-04 (work in progress), March 2018.
- [I-D.ietf-spring-segment-routing-mpls]
Bashandy, A., Filsfils, C., Previdi, S., Decraene, B., Litkowski, S., and R. Shakir, "Segment Routing with MPLS data plane", draft-ietf-spring-segment-routing-mpls-14 (work in progress), June 2018.

- [RFC2104] Krawczyk, H., Bellare, M., and R. Canetti, "HMAC: Keyed-Hashing for Message Authentication", RFC 2104, DOI 10.17487/RFC2104, February 1997, <<https://www.rfc-editor.org/info/rfc2104>>.
- [RFC2827] Ferguson, P. and D. Senie, "Network Ingress Filtering: Defeating Denial of Service Attacks which employ IP Source Address Spoofing", BCP 38, RFC 2827, DOI 10.17487/RFC2827, May 2000, <<https://www.rfc-editor.org/info/rfc2827>>.
- [RFC3032] Rosen, E., Tappan, D., Fedorkow, G., Rekhter, Y., Farinacci, D., Li, T., and A. Conta, "MPLS Label Stack Encoding", RFC 3032, DOI 10.17487/RFC3032, January 2001, <<https://www.rfc-editor.org/info/rfc3032>>.
- [RFC4302] Kent, S., "IP Authentication Header", RFC 4302, DOI 10.17487/RFC4302, December 2005, <<https://www.rfc-editor.org/info/rfc4302>>.
- [RFC4942] Davies, E., Krishnan, S., and P. Savola, "IPv6 Transition/Co-existence Security Considerations", RFC 4942, DOI 10.17487/RFC4942, September 2007, <<https://www.rfc-editor.org/info/rfc4942>>.
- [RFC5308] Hopps, C., "Routing IPv6 with IS-IS", RFC 5308, DOI 10.17487/RFC5308, October 2008, <<https://www.rfc-editor.org/info/rfc5308>>.
- [RFC5340] Coltun, R., Ferguson, D., Moy, J., and A. Lindem, "OSPF for IPv6", RFC 5340, DOI 10.17487/RFC5340, July 2008, <<https://www.rfc-editor.org/info/rfc5340>>.
- [RFC6438] Carpenter, B. and S. Amante, "Using the IPv6 Flow Label for Equal Cost Multipath Routing and Link Aggregation in Tunnels", RFC 6438, DOI 10.17487/RFC6438, November 2011, <<https://www.rfc-editor.org/info/rfc6438>>.
- ▲ [RFC6554] Hui, J., Vasseur, JP., Culler, D., and V. Manral, "An IPv6 Routing Header for Source Routes with the Routing Protocol for Low-Power and Lossy Networks (RPL)", RFC 6554, DOI 10.17487/RFC6554, March 2012, <<https://www.rfc-editor.org/info/rfc6554>>.
- [RFC8126] Cotton, M., Leiba, B., and T. Narten, "Guidelines for Writing an IANA Considerations Section in RFCs", BCP 26, RFC 8126, DOI 10.17487/RFC8126, June 2017, <<https://www.rfc-editor.org/info/rfc8126>>.

Mis en forme : Anglais (États Unis)

Authors' Addresses

Clarence Filsfils (editor)
Cisco Systems, Inc.
Brussels
BE

Email: cfilsfil@cisco.com

Stefano Previdi
Individual
Italy

Email: stefano@previdi.net

John Leddy
Comcast
4100 East Dry Creek Road
Centennial, CO 80122
US

Email: John_Leddy@comcast.com

Satoru Matsushima
Softbank

Email: satoru.matsushima@g.softbank.co.jp

Daniel Voyer (editor)
Bell Canada

Email: daniel.voyer@bell.ca