

DOTS  
Internet-Draft  
Intended status: Standards Track  
Expires: June 22, 24, 2018

T. Reddy  
McAfee  
M. Boucadair  
Orange  
P. Patil  
Cisco  
A. Mortensen  
Arbor Networks, Inc.  
N. Teague  
Verisign, Inc.  
December 19, 21, 2017

Distributed Denial-of-Service Open Threat Signaling (DOTS) Signal  
Channel  
~~draft-ietf-dots-signal-channel-14~~  
draft-ietf-dots-signal-channel-15

Abstract

This document specifies the DOTS signal channel, a protocol for signaling the need for protection against Distributed Denial-of-Service (DDoS) attacks to a server capable of enabling network traffic mitigation on behalf of the requesting client.

A companion document defines the DOTS data channel, a separate reliable communication layer for DOTS management and configuration purposes.

Editorial Note (To be removed by RFC Editor)

Please update these statements with the RFC number to be assigned to this document:

- o "This version of this YANG module is part of RFC XXXX;"
- o "RFC XXXX: Distributed Denial-of-Service Open Threat Signaling (DOTS) Signal Channel";
- o "| 3.00 | Alternate server | [RFCXXXX] |"
- o reference: RFC XXXX
- o This RFC

Please update TBD statements with the port number to be assigned to DOTS Signal Channel Protocol.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on June 22, 24, 2018.

Copyright Notice

Copyright (c) 2017 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Simplified BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Simplified BSD License.

Table of Contents

1. Introduction . . . . .	3
---------------------------	---

2. Notational Conventions and Terminology . . . . .	5	
3. Design Overview . . . . .	6	
4. DOTS Signal Channel: Messages & Behaviors . . . . .	8	
4.1. DOTS Server(s) Discovery . . . . .	8	
4.2. CoAP URIs . . . . .	8	
4.3. Happy Eyeballs for DOTS Signal Channel . . . . .	9	
4.4. DOTS Mitigation Methods . . . . .	10	
4.4.1. Request Mitigation . . . . .	11	
4.4.2. Retrieve Information Related to a Mitigation . . . . .	20	22
4.4.3. Efficacy Update from DOTS Clients . . . . .	28	31
4.4.4. Withdraw a Mitigation . . . . .	30	33
4.5. DOTS Signal Channel Session Configuration . . . . .	32	35
4.5.1. Discover Configuration Parameters . . . . .	33	36
4.5.2. Convey DOTS Signal Channel Session Configuration . . . . .	37	40
4.5.3. Delete DOTS Signal Channel Session Configuration . . . . .	43	46
4.6. Redirected Signaling . . . . .	44	47
4.7. Heartbeat Mechanism . . . . .	45	48
5. DOTS Signal Channel YANG Module . . . . .	47	50
5.1. Tree Structure . . . . .	47	50
5.2. YANG Module . . . . .	49	52
6. Mapping Parameters to CBOR . . . . .	62	65
7. (D)TLS Protocol Profile and Performance Considerations . . . . .	63	66
7.1. (D)TLS Protocol Profile . . . . .	63	66
7.2. (D)TLS 1.3 Considerations . . . . .	64	68
7.3. MTU and Fragmentation . . . . .	65	69
8. Mutual Authentication of DOTS Agents & Authorization of DOTS Clients . . . . .	66	70
9. IANA Considerations . . . . .	68	71
9.1. DOTS Signal Channel UDP and TCP Port Number . . . . .	68	71
9.2. Well-Known 'dots' URI . . . . .	68	71
9.3. CoAP Response Code . . . . .	68	71
9.4. DOTS Signal Channel CBOR Mappings Registry . . . . .	69	72
9.4.1. Registration Template . . . . .	69	72
9.4.2. Initial Registry Contents . . . . .	69	72
9.5. DOTS Signal Channel YANG Module . . . . .	75	78
10. Implementation Status . . . . .	75	79
10.1. nttdots . . . . .	76	79
11. Security Considerations . . . . .	76	80
12. Contributors . . . . .	77	80
13. Acknowledgements . . . . .	77	81
14. References . . . . .	78	81
14.1. Normative References . . . . .	78	81
14.2. Informative References . . . . .	80	83
Authors' Addresses . . . . .	84	87

## 1. Introduction

A distributed denial-of-service (DDoS) attack is an attempt to make machines or network resources unavailable to their intended users. In most cases, sufficient scale can be achieved by compromising enough end-hosts and using those infected hosts to perpetrate and amplify the attack. The victim in this attack can be an application server, a host, a router, a firewall, or an entire network.

Network applications have finite resources like CPU cycles, the number of processes or threads they can create and use, the maximum number of simultaneous connections it can handle, the limited resources of the control plane, etc. When processing network traffic, such applications are supposed to use these resources to offer the intended task in the most efficient manner. However, a DDoS attacker may be able to prevent an application from performing its intended task by making the application exhaust its finite resources.

TCP DDoS SYN-flood, for example, is a memory-exhausting attack while ACK-flood is a CPU-exhausting attack [RFC4987]. Attacks on the link are carried out by sending enough traffic so that the link becomes congested, thereby likely causing packet loss for legitimate traffic. Stateful firewalls can also be attacked by sending traffic that causes the firewall to maintain an excessive number of states that may jeopardize the firewall's operation overall, besides like performance impacts. The firewall then runs out of memory, and can no longer instantiate the states required to process legitimate flows. Other possible DDoS attacks are discussed in [RFC4732].

In many cases, it may not be possible for network administrators to determine the cause(s) of an attack. They may instead just realize that certain resources seem to be under attack. This document defines a lightweight protocol that allows a DOTS client to request mitigation from one or more DOTS servers for protection against detected, suspected, or anticipated attacks. This protocol enables cooperation between DOTS agents to permit a highly-automated network defense that is robust, reliable, and secure.

An example of a network diagram that illustrates a deployment of DOTS agents is shown in Figure 1. In this example, a DOTS server is operating on the access network. A DOTS client is located on the LAN (Local Area Network), while a DOTS gateway is embedded in the CPE (Customer Premises Equipment).




Figure 1: Sample DOTS Deployment (1)

DOTS servers can also be reachable over the Internet, as depicted in Figure 2.




Figure 2: Sample DOTS Deployment (2)

In typical deployments, the DOTS client belongs to a different administrative domain than the DOTS server. For example, the DOTS client is embedded in a firewall protecting services owned and operated by a domain, while the DOTS server is owned and operated by a different domain providing DDoS mitigation services. The latter might or might not provide connectivity services to the network hosting the DOTS client.

The DOTS server may (not) be co-located with the DOTS mitigator. In typical deployments, the DOTS server belongs to the same administrative domain as the mitigator. The DOTS client can communicate directly with a DOTS server or indirectly via a DOTS gateway.

The document adheres to the DOTS architecture [I-D.ietf-dots-architecture]. The requirements for DOTS signal channel protocol are documented in [I-D.ietf-dots-requirements]. This document satisfies all the use cases discussed in [I-D.ietf-dots-use-cases].

This document focuses on the DOTS signal channel. This is a companion document of the DOTS data channel specification [I-D.ietf-dots-data-channel] that defines a configuration and a bulk data exchange mechanism supporting the DOTS signal channel.

## 2. Notational Conventions and Terminology

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC2119].

(D)TLS is used for statements that apply to both Transport Layer Security [RFC5246] and Datagram Transport Layer Security [RFC6347]. Specific terms are used for any statement that applies to either protocol alone.

The reader should be familiar with the terms defined in [I-D.ietf-dots-architecture].

The meaning of the symbols in YANG tree diagrams is defined in [I-D.ietf-netmod-yang-tree-diagrams].

## 3. Design Overview

The DOTS signal channel is built on top of the Constrained Application Protocol (CoAP) [RFC7252], a lightweight protocol originally designed for constrained devices and networks. The many features of CoAP (expectation of packet loss, support for asynchronous non-confirmable messaging, congestion control, small message overhead limiting the need for fragmentation, use of minimal resources, and support for (D)TLS) makes it a good candidate to build the DOTS signaling mechanism from.

The DOTS signal channel is layered on existing standards (Figure 3).




Figure 3: Abstract Layering of DOTS signal channel over CoAP over (D)TLS

By default, a DOTS signal channel MUST run over port number TBD as defined in Section 9.1, for both UDP and TCP, unless the DOTS server has a mutual agreement with its DOTS clients to use a different port number. DOTS clients may alternatively support means to dynamically discover the ports used by their DOTS servers. In order to use a distinct port number (as opposed to TBD), DOTS clients and servers should support a configurable parameter to supply the port number to use. The rationale for not using the default port number 5684 ((D)TLS CoAP) is to allow for differentiated behaviors in environments where both a DOTS gateway and an IoT gateway (e.g., Figure 3 of [RFC7452]) are present.

The signal channel is initiated by the DOTS client (Section 4.4). Once the signal channel is established, the DOTS agents periodically send heartbeats to keep the channel active (Section 4.7). At any time, the DOTS client may send a mitigation request message to a DOTS server over the active channel. While mitigation is active because of the higher likelihood of packet loss during a DDoS attack, the DOTS server periodically sends status messages to the client, including basic mitigation feedback details. Mitigation remains active until the DOTS client explicitly terminates mitigation, or the mitigation lifetime expires.

DOTS signaling can happen with DTLS [RFC6347] over UDP and TLS [RFC5246] over TCP. Likewise, DOTS requests may be sent using IPv4 or IPv6 transfer capabilities. A Happy Eyeballs procedure for DOTS signal channel is specified in Section 4.3.

Messages exchanged between DOTS agents are serialized using Concise Binary Object Representation (CBOR) [RFC7049], CBOR is a binary encoding scheme designed for small code and message size. CBOR-encoded payloads are used to carry signal channel-specific payload messages which convey request parameters and response information such as errors. In order to allow the use of the same data models, [RFC7951] specifies the JSON encoding of YANG-modeled data. A similar effort for CBOR is defined in [I-D.ietf-core-yang-cbor]. All parameters in the payload of the DOTS signal channel are mapped to CBOR types as specified in Section 6.

From that standpoint, this document specifies a YANG data model for representing mitigation scopes and DOTS signal channel session configuration data (Section 5). Representing these data as CBOR data is assumed to follow the rules in [I-D.ietf-core-yang-cbor] or those in [RFC7951] combined with JSON/CBOR conversion rules in [RFC7049].

In order to prevent fragmentation, DOTS agents must follow the recommendations documented in Section 4.6 of [RFC7252]. Refer to Section 7.3 for more details.

DOTS agents MUST support GET, PUT, and DELETE CoAP methods. The payload included in CoAP responses with 2.xx and 3.xx Response Codes MUST be of content type "application/cbor" (Section 5.5.1 of [RFC7252]). CoAP responses with 4.xx and 5.xx error Response Codes MUST include a diagnostic payload (Section 5.5.2 of [RFC7252]). The Diagnostic Payload may contain additional information to aid troubleshooting.

In deployments where multiple DOTS clients are enabled in a network (owned and operated by the same entity), the DOTS server may detect conflicting mitigation requests from these clients. This document does not aim to specify a comprehensive list of conditions under which a DOTS server will characterize two mitigation requests from distinct DOTS clients as conflicting, nor recommend a DOTS server behavior for processing conflicting mitigation requests. Those considerations are implementation- and deployment-specific. Nevertheless, the document specifies the mechanisms to notify DOTS clients when conflicts occur, including the conflict cause (Section 4.4).

In deployments where one or more translators (e.g., Traditional NAT [RFC3022], CGN [RFC6888], NAT64 [RFC6146], NPTv6 [RFC6296]) are enabled between the client's network and the DOTS server, DOTS signal channel messages forwarded to a DOTS server must not include internal IP addresses/prefixes and/or port numbers; external addresses/prefixes and/or port numbers as assigned by the translator must be used instead. This document does not make any recommendation about possible translator discovery mechanisms. The following are some (non-exhaustive) deployment examples that may be considered:

- o Port Control Protocol (PCP) [RFC6887] or Session Traversal Utilities for NAT (STUN) [RFC5389] may be used to retrieve the external addresses/prefixes and/or port numbers. Information retrieved by means of PCP or STUN will be used to feed the DOTS signal channel messages that will be sent to a DOTS server.
- o A DOTS gateway may be co-located with the translator. The DOTS gateway will need to update the DOTS messages, based upon the local translator's binding table.

#### 4. DOTS Signal Channel: Messages & Behaviors

##### 4.1. DOTS Server(s) Discovery

This document assumes that DOTS clients are provisioned with the reachability information of their DOTS server(s) using a variety of means (e.g., local configuration, or dynamic means such as DHCP). These means are out of scope of this document.

Likewise, it is out of scope of this document to specify the behavior of a DOTS client when it sends requests (e.g., contact all servers, select one server among the list) when multiple DOTS servers are provisioned.

##### 4.2. CoAP URIs

The DOTS server MUST support the use of the path-prefix of `"/.well-known/"` as defined in [RFC5785] and the registered name of `"dots"`. Each DOTS operation is indicated by a path-suffix that indicates the intended operation. The operation path (Table 1) is appended to the path-prefix to form the URI used with a CoAP request to perform the desired DOTS operation.

Operation	Operation path	Details
Mitigation	<code>/v1/mitigate</code>	Section 4.4
Session configuration	<code>/v1/config</code>	Section 4.5

Table 1: Operations and their corresponding URIs

##### 4.3. Happy Eyeballs for DOTS Signal Channel

[I-D.ietf-dots-requirements] mentions that DOTS agents will have to support both connectionless and connection-oriented protocols. As such, the DOTS signal channel is designed to operate with DTLS over UDP and TLS over TCP. Further, a DOTS client may acquire a list of IPv4 and IPv6 addresses (Section 4.1), each of which can be used to contact the DOTS server using UDP and TCP. The following specifies the procedure to follow to select the address family and the transport protocol for sending DOTS signal channel messages.

Such procedure is needed to avoid experiencing long connection delays. For example, if an IPv4 path to reach a DOTS server is found, but the DOTS server's IPv6 path is not working, a dual-stack DOTS client may experience a significant connection delay compared to an IPv4-only DOTS client. The other problem is that if a middlebox between the DOTS client and DOTS server is configured to block UDP traffic, the DOTS client will fail to establish a DTLS session with the DOTS server and, as a consequence, will have to fall back to TLS over TCP, thereby incurring significant connection delays.

To overcome these connection setup problems, the DOTS client attempts to connect to its DOTS server(s) using both IPv6 and IPv4, and tries both DTLS over UDP and TLS over TCP in a manner similar to the Happy Eyeballs mechanism [RFC6555]. These connection attempts are performed by the DOTS client when it initializes. The results of the Happy Eyeballs procedure are used by the DOTS client for sending its subsequent messages to the DOTS server.

The order of preference of DOTS signal channel address family and

transport protocol (most preferred first) is: UDP over IPv6, UDP over IPv4, TCP over IPv6, and finally TCP over IPv4. This order adheres to the address preference order specified in [RFC6724] and the DOTS signal channel preference which privileges the use of UDP over TCP (to avoid TCP's head of line blocking).

In reference to Figure 4, the DOTS client sends two TCP SYNs and two DTLS ClientHello messages at the same time over IPv6 and IPv4. In this example, it is assumed that the IPv6 path is broken and UDP traffic is dropped by a middlebox but has little impact to the DOTS client because there is no long delay before using IPv4 and TCP. The DOTS client repeats the mechanism to discover whether DOTS signal channel messages with DTLS over UDP becomes available from the DOTS server, so the DOTS client can migrate the DOTS signal channel from TCP to UDP. Such probing SHOULD NOT be done more frequently than every 24 hours and MUST NOT be done more frequently than every 5 minutes.

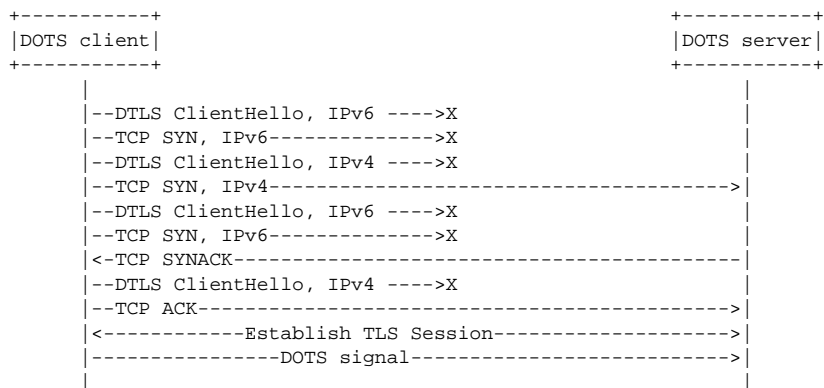


Figure 4: DOTS Happy Eyeballs

#### 4.4. DOTS Mitigation Methods

The following methods are used by a DOTS client to request, withdraw, or retrieve the status of mitigation requests:

**PUT:** DOTS clients use the PUT method to request mitigation from a DOTS server (Section 4.4.1). During active mitigation, DOTS clients may use PUT requests to carry mitigation efficacy updates to the DOTS server (Section 4.4.3).

**GET:** DOTS clients may use the GET method to subscribe to DOTS server status messages, or to retrieve the list of its mitigations maintained by a DOTS server (Section 4.4.2).

**DELETE:** DOTS clients use the DELETE method to withdraw a request for mitigation from a DOTS server (Section 4.4.4).

Mitigation request and response messages are marked as Non-confirmable messages (Section 2.2 of [RFC7252]).

DOTS agents SHOULD follow the data transmission guidelines discussed in Section 3.1.3 of [RFC8085] and control transmission behavior by not sending more than one UDP datagram per RTT to the peer DOTS agent on average.

Requests marked by the DOTS client as Non-confirmable messages are sent at regular intervals until a response is received from the DOTS server. If the DOTS client cannot maintain an RTT estimate, it SHOULD NOT send more than one Non-confirmable request every 3 seconds, and SHOULD use an even less aggressive rate whenever possible (case 2 in Section 3.1.3 of [RFC8085]).

##### 4.4.1. Request Mitigation

When a DOTS client requires mitigation for some reason, the DOTS client uses the CoAP PUT method to send a mitigation request to its DOTS server(s) (Figure 5, illustrated in JSON diagnostic notation).

If this DOTS client is entitled to solicit the DOTS service, the DOTS server can enable mitigation on behalf of the DOTS client by communicating the DOTS client's request to the mitigator and relaying selected mitigator feedback to the requesting DOTS client.

```

Header: PUT (Code=0.03)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
  
```

```

Uri-Path: "version"
Uri-Path: "mitigate"
Content-Type: "application/cbor"
{
  "mitigation-scope": {
    "client-identifier":
    "client-domain-hash": [
      "string"
    ],
    "request-nonce": "string",
    "scope": [
      {
        "mitigation-id": integer,
        "target-prefix": [
          "string"
        ],
        "target-port-range": [
          {
            "lower-port": integer,
            "upper-port": integer
          }
        ],
        "target-protocol": [
          integer
        ],
        "target-fqdn": [
          "string"
        ],
        "target-uri": [
          "string"
        ],
        "alias-name": [
          "string"
        ],
        "lifetime": integer
      }
    ]
  }
}

```

Figure 5: PUT to convey DOTS mitigation requests

The parameters are described below:

~~client-identifier:~~ The client

request-nonce: A globally unique identifier **MAY** that is meant to prevent collisions among DOTS clients. It **MUST** be conveyed generated by a ~~server-side~~ DOTS gateway clients. A variety of methods can be used to ~~propagate the DOTS client identity from~~ generate such identifier, e.g., cryptographic means [RFC4086], mimic the ~~gateway's~~ algorithm in [RFC4941], etc.

'request-nonce' is particularly useful for client-side DOTS gateways to ~~the gateway's server-side, and~~

~~from the gateway's server-side~~ uniquely identify requests relayed to ~~the~~ an upstream DOTS server. ~~'client-identifier'~~ **MAY** be used by

DOTS servers **MUST** echo the ~~final~~ value of 'request-nonce' in the responses sent to DOTS ~~server~~ clients.

This is a mandatory attribute.

mitigation-id: Identifier for **policy** ~~enforcement purposes.~~

~~The 'client-identifier' value the mitigation request represented with an integer. This identifier MUST be assigned by unique for each mitigation request bound to the server-side DOTS gateway in a manner that ensures that there is zero probability that client, i.e., the same 'mitigation-id' parameter value will in the mitigation request needs to be assigned unique relative to a different DOTS client. The server-side DOTS gateway MUST conceal potentially sensitive DOTS client identity information.~~

~~If aggregating DOTS the 'mitigation-id' parameter values of active mitigation requests received conveyed from multiple DOTS clients is enabled, the server-side DOTS gateway has to include a list of 'client-identifier' values; each value is pointing to a unique DOTS client that is in the aggregated list. It is out of scope of this document to specify how aggregation is implemented by a the~~

DOTS ~~gateway.~~

~~The client identifier attribute MUST NOT be generated and included by DOTS clients.~~

~~DOTS servers MUST ignore client identifier attributes that are directly supplied by source DOTS clients. This implies that first server-side DOTS gateways MUST strip client identifier attributes supplied by DOTS clients. DOTS servers MAY support a configuration parameter to identify DOTS gateways that are trusted to supply client identifier attributes.~~

~~This is an optional attribute.~~

~~mitigation-id: Identifier for the mitigation request represented with an integer. This identifier MUST be unique for each mitigation request bound to the DOTS client, i.e., the 'mitigation-id' parameter value in the mitigation request needs to be unique relative to the 'mitigation-id' parameter values of active mitigation requests conveyed from the DOTS client to the DOTS server. This identifier~~ **server. This identifier** MUST be generated by the DOTS client. This document does not make any assumption about how this identifier is generated.

This is a mandatory attribute.

target-prefix: A list of prefixes identifying resources under attack. Prefixes are represented using Classless Inter-Domain Routing (CIDR) notation [RFC4632].  
As a reminder, the prefix length must be less than or equal to 32 (resp. 128) for IPv4 (resp. IPv6).

This is an optional attribute.

target-port-range: A list of port numbers bound to resources under attack.

The port range is defined by two bounds, a lower port number (lower-port) and an upper port number (upper-port). When only 'lower-port' is present, it represents a single port number. For TCP, UDP, Stream Control Transmission Protocol (SCTP) [RFC4960], or Datagram Congestion Control Protocol (DCCP) [RFC4340], the range of ports can be, for example, 1024-65535.

This is an optional attribute.

target-protocol: A list of protocols involved in an attack. Values are taken from the IANA protocol registry [proto\_numbers].

The value '0' has a special meaning for 'all protocols'.

This is an optional attribute.

target-fqdn: A list of Fully Qualified Domain Names (FQDNs) identifying resources under attack. An FQDN is the full name of a resource, rather than just its hostname. For example, "venera" is a hostname, and "venera.isi.edu" is an FQDN.

This is an optional attribute.

target-uri: A list of Uniform Resource Identifiers (URIs) [RFC3986] identifying resources under attack.

This is an optional attribute.

alias-name: A list of aliases of resources for which the mitigation is requested. Aliases can be created using the DOTS data channel (Section 6.1 of [I-D.ietf-dots-data-channel]), direct configuration, or other means. An alias is used in subsequent signal channel exchanges to refer more efficiently to the resources under attack.

This is an optional attribute.

lifetime: Lifetime of the mitigation request in seconds. The RECOMMENDED lifetime of a mitigation request is 3600 seconds (60 minutes) -- this value was chosen to be long enough so that refreshing is not typically a burden on the DOTS client, while expiring the request where the client has unexpectedly quit in a timely manner. DOTS clients MUST include this parameter in their mitigation requests. Upon the expiry of this lifetime, and if the request is not refreshed, the mitigation request is removed. The request can be refreshed by sending the same request again.



A lifetime of 0 '0' in a mitigation request is an invalid value.

A lifetime of negative one (-1) indicates indefinite lifetime for the mitigation request. The DOTS server MAY refuse indefinite lifetime, for policy reasons; the granted lifetime value is returned in the response. DOTS clients MUST be prepared to not be granted mitigations with indefinite lifetimes.

The DOTS server MUST always indicate the actual lifetime in the response and the remaining lifetime in status messages sent to the the DOTS client.

This is a mandatory attribute.

In deployments where server-side DOTS gateways are enabled, some information about the origin source client domain has to be supplied to the DOTS server. That information is meant to assist the DOTS server to enforce some policies. Figure 6 shows an example of a request relayed by a server-side DOTS gateway.

```
Header: PUT (Code=0.03)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "mitigate"
Content-Type: "application/cbor"
{
  "mitigation-scope": {
    "client-domain-hash": [
      "string"
    ],
    "request-nonce": "string",
    "scope": [
      {
        "mitigation-id": integer,
        "target-prefix": [
          "string"
        ],
        "target-port-range": [
          {
            "lower-port": integer,
            "upper-port": integer
          }
        ],
        "target-protocol": [
          integer
        ],
        "target-fqdn": [
          "string"
        ],
        "target-uri": [
          "string"
        ],
        "alias-name": [
          "string"
        ],
        "lifetime": integer
      }
    ]
  }
}
```

Figure 6: PUT to convey DOTS mitigation request as relayed by a server-side DOTS gateway

The DOTS gateway may add the following parameter:

**client-domain-hash:** The client identifier MAY be conveyed by a server-side DOTS gateway to propagate the source domain identity from the gateway's client-side to the gateway's server-side, and from the gateway's server-side to the DOTS server. 'client-domain-hash' MAY be used by the final DOTS server for policy enforcement purposes (e.g., enforce a quota on filtering rules).

The 'client-domain-hash' value MUST be assigned by the server-side DOTS gateway in a manner that ensures that there is zero probability that the same value will be assigned to a different client domain. The server-side DOTS gateway MUST conceal potentially sensitive DOTS client identity information.

If the DOTS client is using the certificate provisioned by the Enrollment over Secure Transport (EST) server [RFC7030] in the DOTS gateway-domain to authenticate itself to the DOTS gateway,

the 'client-domain-hash' value may be the output of a cryptographic hash algorithm whose input is the DER-encoded ASN.1 representation of the Subject Public Key Info (SPKI) of an X.509 certificate. In this version of the specification, the cryptographic hash algorithm used is SHA-256 [RFC6234]. The output of the cryptographic hash algorithm is truncated to 16 bytes; truncation is done by stripping off the final 16 bytes. The truncated output is base64url encoded.

The 'client-domain-hash' attribute **MUST NOT** be generated and included by DOTS clients.

DOTS servers **MUST** ignore 'client-domain-hash' attributes that are directly supplied by source DOTS clients or client-side DOTS gateways. This implies that first server-side DOTS gateways **MUST** strip 'client-domain-hash' attributes supplied by DOTS clients. DOTS servers **MAY** support a configuration parameter to identify DOTS ~~client-~~ gateways that are trusted to supply 'client-domain-hash' attributes.

Only single-valued 'client-domain-hash' are defined in this document.

This is ~~a mandatory~~ **an optional** attribute.

Because of the complexity to handle partial failure cases, this specification does not allow for including multiple mitigation requests in the same PUT request. Concretely, a DOTS client **MUST NOT** include multiple 'scope' parameters in the same PUT request.

The CBOR key values for the parameters are defined in Section 6. Section 9 defines how the CBOR key values can be allocated to standard bodies and vendors.

FQDN and URI mitigation scopes may be thought of as a form of scope alias, in which the addresses to which the domain name or URI resolve represent the full scope of the mitigation.

In the PUT request at least one of the attributes 'target-prefix' or 'target-fqdn' or 'target-uri' or 'alias-name' **MUST** be present. Attributes with ~~empty~~ **empty** values **MUST NOT** be present in a request.

The relative order of two mitigation requests from a DOTS client is determined by comparing their respective 'mitigation-id' values. If two mitigation requests have overlapping mitigation scopes, the mitigation request with the highest numeric 'mitigation-id' value will override the other mitigation request. Two mitigation-ids from a DOTS client are overlapping if there is a common IP address, IP prefix, FQDN, URI, or alias-name. To avoid maintaining a long list of overlapping mitigation requests from a DOTS client and avoid error-prone provisioning of mitigation requests from a DOTS client, the overlapped lower numeric 'mitigation-id' **MUST** be automatically deleted and no longer available at the DOTS server.

The Uri-Path option carries a major and minor version nomenclature to manage versioning and DOTS signal channel in this specification uses v1 major version.

Figure 6 7 shows a PUT request example to signal that ports 80, 8080, and 443 used by 2001:db8:6401::1 and 2001:db8:6401::2 servers are under attack (illustrated in JSON diagnostic notation). **The presence of 'client-domain-hash' indicates that a server-side DOTS gateway has modified the initial PUT request sent by the DOTS client.**

```
Header: PUT (Code=0.03)
Uri-Host: "www.example.com"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "v1"
Uri-Path: "mitigate"
Content-Format: "application/cbor"
{
  "mitigation-scope": {
    "client-identifier":
    "client-domain-hash": [
      "dz6pHjaADkaFTbjr0JGBpw"
    ],
    "request-nonce": "8UfNmBc8lShRkQkmBfJO",
    "scope": [
      {
        "mitigation-id": 12332,
        "target-prefix": [
          "2001:db8:6401::1/128",
          "2001:db8:6401::2/128"
        ]
      }
    ]
  }
}
```

```

    ],
    "target-port-range": [
      {
        "lower-port": 80
      },
      {
        "lower-port": 443
      },
      {
        "lower-port": 8080
      }
    ],
    "target-protocol": [
      6
    ]
  ]
}
}
}

```

Figure 6+ 7: PUT for DOTS ~~signal~~ mitigation request

The corresponding CBOR encoding format is shown in Figure 7- 8.

```

A1                                     # map(1)
01                                     # unsigned(1)
A2                                     # map(2)
  18 20                               # unsigned(32)
  81                                   # array(1)
  76                                   # text(22)
    647A3670486A6141446B614654626A72304A47427077 # "dz6pHjaADkaFTbjr0JGBpw"
  74                                   # text(20)
    3855664E6D6263386C5368526B516B6D42664A4F # "8UfNmBc8lShRkQkmBfJO"
  02                                   # unsigned(2)
  81                                   # array(1)
    A4                                # map(4)
      03                               # unsigned(3)
      19 302C                         # unsigned(12332)
      04                               # unsigned(4)
      82                               # array(2)
        74                             # text(20)
          323030313A6462383A363430313A3A312F313238 # "2001:db8:6401::1/128"
        74                             # text(20)
          323030313A6462383A363430313A3A322F313238 # "2001:db8:6401::2/128"
    05                               # unsigned(5)
    83                               # array(3)
      A1                              # map(1)
        06                             # unsigned(6)
        18 50                         # unsigned(80)
      A1                              # map(1)
        06                             # unsigned(6)
        19 01BB                       # unsigned(443)
      A1                              # map(1)
        06                             # unsigned(6)
        19 1F90                       # unsigned(8080)
    08                               # unsigned(8)
    81                               # array(1)
      06                               # unsigned(6)

```

Figure 7- 8: PUT for DOTS signal (CBOR)

~~If the DOTS client is using the certificate provisioned by the Enrollment over Secure Transport (EST) server [RFC7030] in the DOTS gateway domain to authenticate itself to the DOTS gateway, then the 'client-identifier' value can be the output of a cryptographic hash algorithm whose input is the DER-encoded ASN.1 representation of the Subject Public Key Info (SPKI) of an X.509 certificate.~~

~~In this version of the specification, the cryptographic hash algorithm used is SHA-256 [RFC6234]. The output of the cryptographic hash algorithm is truncated to 16 bytes; truncation is done by stripping off the final 16 bytes. The truncated output is base64url encoded.~~

In both DOTS signal and data channel sessions, the DOTS client MUST authenticate itself to the DOTS server (Section 8). The DOTS server may use the algorithm presented in Section 7 of [RFC7589] to derive the DOTS client identity or username from the client certificate. The DOTS client identity allows the DOTS server to accept mitigation requests with scopes that the DOTS client is authorized to manage. The DOTS server couples the DOTS signal and data channel sessions using the DOTS client identity **and or** the **'client-identifier'** **'client-domain-hash'** parameter value, so the DOTS server can validate whether the aliases conveyed in the mitigation request were indeed created by the same DOTS client

using the DOTS data channel session. If the aliases were not created by the DOTS client, the DOTS server **returns MUST return** 4.00 (Bad Request) in the response.

The DOTS server couples the DOTS signal channel sessions using the DOTS client identity **and or** the ~~'client-identifier'~~ **'client-domain-hash'** parameter value, and the DOTS server uses 'mitigation-id' **and 'request-nonce'** parameter **value values** to detect duplicate mitigation requests. If the mitigation request contains the alias-name and other parameters identifying the target resources (such as, 'target-prefix', 'target-port-range', 'target-fqdn', or 'target-uri'), **then** the DOTS server appends the parameter values in 'alias-name' with the corresponding parameter values in ~~'target-~~~~prefix'~~, **'target-prefix'**, 'target-port-range', 'target-fqdn', or 'target-uri'.

The DOTS server indicates the result of processing the PUT request using CoAP response codes. CoAP 2.xx codes are success. CoAP 4.xx codes are some sort of invalid requests (client errors). CoAP 5.xx codes are returned if the DOTS server has erred or is currently unavailable to provide mitigation in response to the mitigation request from the DOTS client.

Figure 8 9 shows an example of a PUT request that is successfully processed (i.e., CoAP 2.xx response codes).

```
{
  "mitigation-scope": {
    "client-identifier":
    "client-domain-hash": [
      "string"
    ],
    "request-nonce": "string",
    "scope": [
      {
        "mitigation-id": 12332,
        "lifetime": 3600
      }
    ]
  }
}
```

Figure ~~8+~~ 9: 2.xx response body

If the request is missing one or more mandatory attributes, or includes multiple 'scope' parameters, or contains invalid or unknown parameters, the DOTS server MUST reply with 4.00 (Bad Request). DOTS agents can safely ignore Vendor-Specific parameters they don't understand.

A DOTS server that receives a mitigation request with a lifetime set to '0' MUST reply with a 4.00 (Bad Request).

If the DOTS server does not find the 'mitigation-id' parameter value conveyed in the PUT request in its configuration data, it MAY accept the mitigation request by sending back a 2.01 (Created) response to the DOTS client; the DOTS server will consequently try to mitigate the attack.

If the DOTS server finds the 'mitigation-id' parameter value conveyed in the PUT request in its configuration ~~data,~~ **data bound to that DOTS client**, it MAY update the mitigation request, and a 2.04 (Changed) response is returned to indicate a successful update of the mitigation request.

If the request is conflicting with an existing mitigation request from a different DOTS client, and the DOTS server decides to maintain the conflicting mitigation request, the DOTS server returns 4.09 (Conflict) [RFC8132] to the requesting DOTS client. The response includes enough information for a DOTS client to recognize the source of the conflict (refer to 'conflict-information' specified in Section 4.4.2).

For a mitigation request to continue beyond the initial negotiated lifetime, the DOTS client has to refresh the current mitigation request by sending a new PUT request. This PUT request MUST use the same 'mitigation-id' value, and MUST repeat all the other parameters as sent in the original mitigation request apart from a possible change to the lifetime parameter ~~value-~~ **value and 'nonce-request'**.

The DOTS gateway, which inserted a ~~'client-identifier'~~ **'client-domain-hash'** attribute in a request, MUST strip the ~~'client-identifier'~~ **'client-domain-hash'** parameter in the corresponding response before forwarding the response to the DOTS client.

#### 4.4.2. Retrieve Information Related to a Mitigation

A GET request is used by a DOTS client to retrieve information (including status) of DOTS mitigations from a DOTS server.

The same considerations for manipulating ~~'client-identifier'~~ ~~'client-domain-hash'~~ parameter by a **server-side** DOTS **gateway** **gateways** specified in Section 4.4.1 MUST be followed for GET requests.

If the DOTS server does not find the 'mitigation-id' parameter value conveyed in the GET request in its configuration data for the requesting DOTS client or the one identified by ~~'client-identifier'~~, ~~'client-domain-hash'~~, it MUST respond with a 4.04 (Not Found) error response code.

Likewise, the same error MUST be returned as a response to a request to retrieve all mitigation records of a given DOTS client if the DOTS server does not find any mitigation record for that DOTS client or the one identified by ~~'client-identifier'~~, ~~'client-domain-hash'~~.

The 'c' (content) parameter and its permitted values defined in [I-D.ietf-core-comi] can be used to retrieve non-configuration data (attack mitigation status) or configuration data or both. The DOTS server may support this optional filtering capability. It can safely ignore it if not supported.

The following examples illustrate how a DOTS client retrieves active mitigation requests from a DOTS server. In particular:

- o Figure 9 10 shows the example of a GET request to retrieve all DOTS mitigation requests signaled by a DOTS client.
- o Figure 10 11 shows the example of a GET request to retrieve a specific DOTS mitigation request signaled by a DOTS client. The configuration data to be reported in the response is formatted in the same order it was processed by the DOTS server.

These two examples assume the default of "c=a"; that is, the DOTS client asks for all data to be reported by the DOTS server.

```
Header: GET (Code=0.01)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "mitigate"
Observe : 0
{
  "mitigation-scope": {
    "client-identifier": {
    "dz6pHjaADkaFTbjr0JGBpw"
  }
  "request-nonce": "wLjlAepU50LzYICsXr13",
}
```

Figure 9+ 10: GET to retrieve all DOTS mitigation requests

```
Header: GET (Code=0.01)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "mitigate"
Observe : 0
Content-Format: "application/cbor"
{
  "mitigation-scope": {
    "client-identifier": {
    "dz6pHjaADkaFTbjr0JGBpw"
  },
  "request-nonce": "wLjlAepU50LzYICsXr13",
  "scope": [
    {
      "mitigation-id": 12332
    }
  ]
}
```

Figure 10+ 11: GET to retrieve a specific DOTS mitigation request

Figure 11 12 shows a response example of all active mitigation requests associated with the DOTS client on the DOTS server and the mitigation

status of each mitigation request.

```
{
  "mitigation-scope": {
    "request-nonce": "wLjlAepU5OLzYICsXr13",
    "scope": [
      {
        "mitigation-id": 12332,
        "mitigation-start": 1507818434.00, 1507818434,
        "target-protocol": [
          17
        ],
        "lifetime": 1800,
        "status": 2,
        "bytes-dropped": 134334555,
        "bps-dropped": 43344,
        "pkts-dropped": 333334444,
        "pps-dropped": 432432
      },
      {
        "mitigation-id": 12333,
        "mitigation-start": 1507818393.00, 1507818393,
        "target-protocol": [
          6
        ],
        "lifetime": 1800,
        "status": 3,
        "bytes-dropped": 0,
        "bps-dropped": 0,
        "pkts-dropped": 0,
        "pps-dropped": 0
      }
    ]
  }
}
```

Figure 11- 12: Response body

The mitigation status parameters are described below:

**mitigation-start:** Mitigation start time is expressed in seconds relative to 1970-01-01T00:00Z in UTC time (Section 2.4.1 of [RFC7049]). The encoding is modified so that the leading tag 1 (epoch-based date/time) MUST be omitted.

This is a mandatory attribute.

**lifetime:** The remaining lifetime of the mitigation request, in seconds.

This is a mandatory attribute.

**status:** Status of attack mitigation. The various possible values of 'status' parameter are explained in Table 2.

This is a mandatory attribute.

**conflict-information:** Indicates that a mitigation request is conflicting with another mitigation request(s) from other DOTS client(s). This optional attribute has the following structure:

**conflict-status:** Indicates the status of a conflicting mitigation request. The following values are defined:

- 1: DOTS server has detected conflicting mitigation requests from different DOTS clients. This mitigation request is currently inactive until the conflicts are resolved. Another mitigation request is active.
- 2: DOTS server has detected conflicting mitigation requests from different DOTS clients. This mitigation request is currently active.
- 3: DOTS server has detected conflicting mitigation requests from different DOTS clients. All conflicting mitigation requests are inactive.

**conflict-cause:** Indicates the cause of the conflict. The following values are defined:

- 1: Overlapping targets. 'conflict-scope' provides more details about the conflicting target clauses.
- 2: Conflicts with an existing white list. This code is

returned when the DDoS mitigation detects source addresses/ prefixes in the white-listed ACLs are attacking the target.

**conflict-scope** Indicates the conflict scope. It may include a list of IP addresses, a list of prefixes, a list of port numbers, a list of target protocols, a list of FQDNs, a list of URIs, a list of alias-names, or references to conflicting ACLs.

**retry-timer** Indicates, in seconds, the time after which the DOTS client may re-issue the same request. The DOTS server returns 'retry-timer' only to DOTS client(s) for which a mitigation request is deactivated. Any retransmission of the same mitigation request before the expiry of this timer is likely to be rejected by the DOTS server for the same reasons.

The retry-timer SHOULD be equal to the lifetime of the active mitigation request resulting in the deactivation of the conflicting mitigation request. The lifetime of the deactivated mitigation request will be updated to (retry-timer + 45 seconds), so the DOTS client can refresh the deactivated mitigation request after retry-timer seconds before expiry of lifetime and check if the conflict is resolved.

**bytes-dropped:** The total dropped byte count for the mitigation request since the attack mitigation is triggered. The count wraps around when it reaches the maximum value of unsigned integer.

This is an optional attribute.

**bps-dropped:** The average number of dropped bytes per second for the mitigation request since the attack mitigation is triggered. This SHOULD be a five-minute average.

This is an optional attribute.

**pkts-dropped:** The total number of dropped packet count for the mitigation request since the attack mitigation is triggered.

This is an optional attribute.

**pps-dropped:** The average number of dropped packets per second for the mitigation request since the attack mitigation is triggered. This SHOULD be a five-minute average.

This is an optional attribute.

Parameter value	Description
1	Attack mitigation is in progress (e.g., changing the network path to re-route the inbound traffic to DOTS mitigator).
2	Attack is successfully mitigated (e.g., traffic is redirected to a DDOS mitigator and attack traffic is dropped).
3	Attack has stopped and the DOTS client can withdraw the mitigation request.
4	Attack has exceeded the mitigation provider capability.
5	DOTS client has withdrawn the mitigation request and the mitigation is active but terminating.
6	Attack mitigation is now terminated.
7	Attack mitigation is withdrawn.
8	Attack mitigation is rejected.

Table 2: Values of 'status' parameter

The observe option defined in [RFC7641] extends the CoAP core protocol with a mechanism for a CoAP client to "observe" a resource on a CoAP server: The client retrieves a representation of the resource and requests this representation be updated by the server as long as the client is interested in the resource. A DOTS client conveys the observe option set to '0' in the GET request to receive unsolicited notifications of attack mitigation status from the DOTS server.

Unidirectional notifications within the bidirectional signal channel allows unsolicited message delivery, enabling asynchronous notifications between the agents. Due to the higher likelihood of packet loss during a DDoS attack, the DOTS server periodically sends attack mitigation status to the DOTS client and also notifies the DOTS client whenever the status of the attack mitigation changes. If the DOTS server cannot maintain a RTT estimate, it SHOULD NOT send more than one unsolicited notification every 3 seconds, and SHOULD use an even less aggressive rate whenever possible (case 2 in Section 3.1.3 of [RFC8085]). **The DOTS server MUST use the same request nonce as the one used by the DOTS client to observe a resource.**

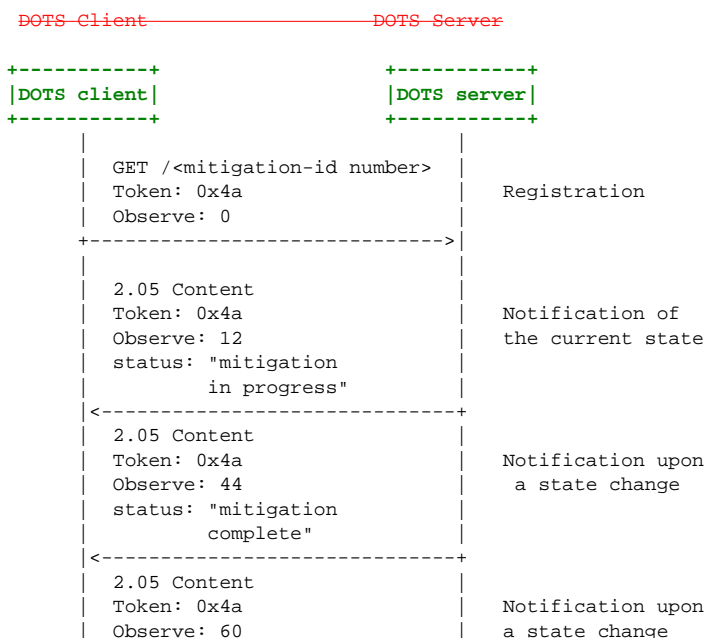
When conflicting requests are detected, the DOTS server enforces the corresponding policy (e.g., accept all requests, reject all requests, accept only one request but reject all the others, ...). It is assumed that this policy is supplied by the DOTS server administrator or it is a default behavior of the DOTS server implementation. Then, the DOTS server sends notification message(s) to the DOTS client(s) at the origin of the conflict. A conflict notification message includes information about the conflict cause, scope, and the status of the mitigation request(s). For example,

- o A notification message with status code set to '8 (Attack mitigation is rejected)' and 'conflict-status' set to '1' is sent to a DOTS client to indicate that this mitigation request is rejected because a conflict is detected.
- o A notification message with status code set to '7 (Attack mitigation is withdrawn)' and 'conflict-status' set to '1' is sent to a DOTS client to indicate that an active mitigation request is deactivated because a conflict is detected.
- o A notification message with status code set to '1 (Attack mitigation is in progress)' and 'conflict-status' set to '2' is sent to a DOTS client to indicate that this mitigation request is in progress, but a conflict is detected.

Upon receipt of a conflict notification message indicating that a mitigation request is deactivated because of a conflict, a DOTS client MUST NOT resend the same mitigation request before the expiry of 'retry-timer'. It is also recommended that DOTS clients support means to alert administrators about mitigation conflicts.

A DOTS client that is no longer interested in receiving notifications from the DOTS server can simply "forget" the observation. When the DOTS server sends the next notification, the DOTS client will not recognize the token in the message and thus will return a Reset message. This causes the DOTS server to remove the associated entry. Alternatively, the DOTS client can explicitly deregister itself by issuing a GET request that has the Token field set to the token of the observation to be cancelled and includes an Observe Option with the value set to '1' (deregister).

Figure 12 13 shows an example of a DOTS client requesting a DOTS server to send notifications related to a given mitigation request.





```

| status: "attack stopped" |
|<-----+
|

```

Figure 12+ 13: Notifications of attack mitigation status

#### 4.4.2.1. Mitigation Status

The DOTS client can send the GET request at frequent intervals without the Observe option to retrieve the configuration data of the mitigation request and non-configuration data (i.e., the attack status). The frequency of polling the DOTS server to get the mitigation status should follow the transmission guidelines given in Section 3.1.3 of [RFC8085]. If the DOTS server has been able to mitigate the attack and the attack has stopped, the DOTS server indicates as such in the status, and the DOTS client recalls the mitigation request by issuing a DELETE request for the ~~mitigation-id~~ 'mitigation-id'.

A DOTS client SHOULD react to the status of the attack as per the information sent by the DOTS server rather than acknowledging by itself, using its own means, that the attack has been mitigated. This ensures that the DOTS client does not recall a mitigation request prematurely because it is possible that the DOTS client does not sense the ~~DDoS~~ DDoS attack on its resources but the DOTS server could be actively mitigating the attack and the attack is not completely averted.

#### 4.4.3. Efficacy Update from DOTS Clients

While DDoS mitigation is active, due to the likelihood of packet loss, a DOTS client MAY periodically transmit DOTS mitigation efficacy updates to the relevant DOTS server. A PUT request is used to convey the mitigation efficacy update to the DOTS server.

The PUT request MUST include all the parameters used in the PUT request to carry the DOTS mitigation request (Section 4.4.1) unchanged apart from the lifetime parameter value. If this is not the case, the DOTS server MUST reject the request with a 4.00 (Bad Request).

The If-Match Option (Section 5.10.8.1 of [RFC7252]) with an empty value is used to make the PUT request conditional on the current existence of the mitigation request. If UDP is used as transport, CoAP requests may arrive out-of-order. For example, the DOTS client may send a PUT request to convey an efficacy update to the DOTS server followed by a DELETE request to withdraw the mitigation request, but the DELETE request arrives at the DOTS server before the PUT request. To handle out-of-order delivery of requests, if an If-Match option is present in the PUT request and the 'mitigation-id' in the request matches a mitigation request from that DOTS client, then the request is processed. If no match is found, the PUT request is silently ignored.

An example of an efficacy update message, which includes an If-Match option with an empty value, is depicted in Figure 13- 14.

```

Header: PUT (Code=0.03)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "mitigate"
Content-Format: "application/cbor"
If-Match:
{
  "mitigation-scope": {
    "client-identifier": {
    "string"
  },
  "request-nonce": "string",
  "scope": [
    {
      "mitigation-id": integer,
      "target-prefix": [
        "string"
      ],
      "target-port-range": [
        {
          "lower-port": integer,
          "upper-port": integer
        }
      ],
      "target-protocol": [

```

```

        integer
      ],
      "target-fqdn": [
        "string"
      ],
      "target-uri": [
        "string"
      ],
      "alias-name": [
        "string"
      ],
      "lifetime": integer,
      "attack-status": integer
    }
  ]
}

```

Figure ~~13~~ 14: Efficacy Update

The 'attack-status' parameter is a mandatory attribute when performing an efficacy update. The various possible values contained in the 'attack-status' parameter are described in Table 3.

Parameter value	Description
1	The DOTS client determines that it is still under attack.
2	The DOTS client determines that the attack is successfully mitigated (e.g., attack traffic is not seen).

Table 3: Values of 'attack-status' parameter

The DOTS server indicates the result of processing a PUT request using CoAP response codes. The response code 2.04 (Changed) is returned if the DOTS server has accepted the mitigation efficacy update. The error response code 5.03 (Service Unavailable) is returned if the DOTS server has erred or is incapable of performing the mitigation.

#### 4.4.4. Withdraw a Mitigation

A DELETE request is used to withdraw a DOTS mitigation request from a DOTS server (Figure ~~14~~ 15).

The same considerations for manipulating ~~'client-identifier'~~ 'client-domain-hash' parameter by a DOTS ~~gateway~~, gateways, as specified in Section 4.4.1, MUST be followed for DELETE requests.

```

Header: DELETE (Code=0.04)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "mitigate"
Content-Format: "application/cbor"
{
  "mitigation-scope": {
    "client-identifier": {
      "string"
    },
    "request-nonce": "string",
    "scope": [
      {
        "mitigation-id": integer
      }
    ]
  }
}

```

Figure ~~14~~ 15: Withdraw DOTS ~~signal~~ mitigation

If the request does not include a 'mitigation-id' parameter, the DOTS server MUST reply with a 4.00 (Bad Request).

Once the request is validated, the DOTS server immediately acknowledges a DOTS client's request to withdraw the DOTS signal using 2.02 (Deleted) response code with no response payload. A 2.02 (Deleted) Response Code is returned even if the 'mitigation-id'

parameter value conveyed in the DELETE request does not exist in its configuration data before the request.

If the DOTS server finds the 'mitigation-id' parameter value conveyed in the DELETE request in its configuration data for the DOTS client, then to protect against route or DNS flapping caused by a DOTS client rapidly removing a mitigation, and to dampen the effect of oscillating attacks, the DOTS server MAY allow mitigation to continue for a limited period after acknowledging a DOTS client's withdrawal of a mitigation request. During this period, the DOTS server status messages SHOULD indicate that mitigation is active but terminating (Section 4.4.2).

The initial active-but-terminating period SHOULD be sufficiently long to absorb latency incurred by route propagation. The active-but-terminating period SHOULD be set by default to 120 seconds. If the client requests mitigation again before the initial active-but-terminating period elapses, the DOTS server MAY exponentially increase the active-but-terminating period up to a maximum of 300 seconds (5 minutes).

After the active-but-terminating period elapses, the DOTS server MUST treat the mitigation as terminated, as the DOTS client is no longer responsible for the mitigation. For example, if there is a financial relationship between the DOTS client and server domains, the DOTS client stops incurring cost at this point.

#### 4.5. DOTS Signal Channel Session Configuration

A DOTS client can negotiate, configure, and retrieve the DOTS signal channel session behavior. The DOTS signal channel can be used, for example, to configure the following:

- a. Heartbeat interval (heartbeat-interval): DOTS agents regularly send heartbeats (CoAP Ping/Pong) to each other after mutual authentication is successfully completed in order to keep the DOTS signal channel open. Heartbeat messages are exchanged between DOTS agents every 'heartbeat-interval' seconds to detect the current status of the DOTS signal channel session.
- b. Missing heartbeats allowed (missing-hb-allowed): This variable indicates the maximum number of consecutive heartbeat messages for which a DOTS agent did not receive a response before concluding that the session is disconnected or defunct.
- c. Acceptable signal loss ratio: Maximum retransmissions, retransmission timeout value, and other message transmission parameters for the DOTS signal channel.

The same or distinct configuration sets may be used during ~~attack~~ (~~'attack-time-config'~~) and ~~peace~~ times (~~'peace-time-config'~~), ~~when a mitigation is active ('mitigating-config')~~ and ~~when no mitigation is active ('idle-config')~~. This is particularly useful for DOTS servers that might want to reduce heartbeat frequency or cease heartbeat exchanges when an active DOTS client has not requested mitigation. If distinct ~~configuration~~ **configurations** are used, DOTS agents MUST follow the appropriate configuration set as a function of the mitigation activity (e.g., if no mitigation request is active, ~~'peace-time-config'-related~~ **'idle-config'-related** values must be followed). Additionally, DOTS agents MUST automatically switch to the other configuration upon a change in the mitigation activity (e.g., if an attack mitigation is launched after a peacetime, the DOTS agent switches from ~~'peace-time-config'~~ **'idle-config'** to ~~'attack-time-config'-related~~ **'mitigating-config'-related** values).

Requests and responses are deemed reliable by marking them as Confirmable (CON) messages. DOTS signal channel session configuration requests and responses are marked as Confirmable messages. As explained in Section 2.1 of [RFC7252], a Confirmable message is retransmitted using a default timeout and exponential back-off between retransmissions, until the DOTS server sends an Acknowledgement message (ACK) with the same Message ID conveyed from the DOTS client.

Message transmission parameters are defined in Section 4.8 of [RFC7252]. The DOTS server can either piggyback the response in the acknowledgement message or, if the DOTS server cannot respond immediately to a request carried in a Confirmable message, it simply responds with an Empty Acknowledgement message so that the DOTS client can stop retransmitting the request. Empty Acknowledgement message is explained in Section 2.2 of [RFC7252]. When the response is ready, the server sends it in a new Confirmable message which in turn needs to be acknowledged by the DOTS client (see Sections 5.2.1 and 5.2.2 of [RFC7252]). Requests and responses exchanged between DOTS agents during peacetime are marked as Confirmable messages.

Implementation Note: A DOTS client that receives a response in a CON message may want to clean up the message state right after sending the ACK. If that ACK is lost and the DOTS server retransmits the CON, the DOTS client may no longer have any state that would help it correlate this response, thereby unexpecting the retransmission message. The DOTS client will send a Reset message so it does not receive any more retransmissions. This behavior is normal and not an indication of an error (see Section 5.3.2 of [RFC7252] for more details).

#### 4.5.1. Discover Configuration Parameters

A GET request is used to obtain acceptable (e.g., minimum and maximum values) and current configuration parameters on the DOTS server for DOTS signal channel session configuration.

Figure 15 16 shows how to obtain acceptable configuration parameters for the DOTS server.

```
Header: GET (Code=0.01)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "config"
Content-Format: "application/cbor"
{
  "signal-config": {
    "session-id": integer,
    "request-nonce": "string"
  }
}
```

Figure 15+ 16: GET to retrieve configuration

The DOTS server in the 2.05 (Content) response conveys the current, minimum, and maximum attribute values acceptable by the DOTS server (Figure 16+ 17).

Content-Format: "application/cbor"

```
{
  "attack-time-config":
  "signal-config": {
    "session-id": integer,
    "request-nonce": "string",
    "mitigating-config": {
      "heartbeat-interval": {
        "current-value": integer,
        "min-value": integer,
        "max-value": integer
      },
      "missing-hb-allowed": {
        "current-value": integer,
        "min-value": integer,
        "max-value": integer
      },
      "max-retransmit": {
        "current-value": integer,
        "min-value": integer,
        "max-value": integer
      },
      "ack-timeout": {
        "current-value": integer,
        "min-value": integer,
        "max-value": integer
      },
      "ack-random-factor": {
        "current-value": number,
        "min-value": number,
        "max-value": number
      }
    }
  },
  "peace-time-config":
  "idle-config": {
    "heartbeat-interval": {
      "current-value": integer,
      "min-value": integer,
      "max-value": integer
    },
    "missing-hb-allowed": {
      "current-value": integer,
      "min-value": integer,
      "max-value": integer
    }
  }
}
```

```

    },
    "max-retransmit": {
      "current-value": integer,
      "min-value": integer,
      "max-value": integer
    },
    "ack-timeout": {
      "current-value": integer,
      "min-value": integer,
      "max-value": integer
    },
    "ack-random-factor": {
      "current-value": number,
      "min-value": number,
      "max-value": number
    }
  },
  "trigger-mitigation": {
    "current-value": boolean
  },
  "config-interval": {
    "current-value": integer,
    "min-value": integer,
    "max-value": integer
  }
}
}

```

Figure 16+ 17: GET response body

Figure 17 18 shows an example of acceptable and current configuration parameters on a DOTS server for DOTS signal channel session configuration. The same acceptable configuration is used during attack and peace times.

Content-Format: "application/cbor"

```

{
  "attack-time-config": {
  "signal-config": {
    "session-id": 589678,
    "request-nonce": "1PfX4iwTQwCuUg0ka3yk",
    "mitigating-config": {
      "heartbeat-interval": {
        "current-value": 30,
        "min-value": 15,
        "max-value": 240
      },
      "missing-hb-allowed": {
        "current-value": 5,
        "min-value": 3,
        "max-value": 9
      },
      "max-retransmit": {
        "current-value": 3,
        "min-value": 2,
        "max-value": 15
      },
      "ack-timeout": {
        "current-value": 2,
        "min-value": 1,
        "max-value": 30
      },
      "ack-random-factor": {
        "current-value": 1.5,
        "min-value": 1.1,
        "max-value": 4.0
      }
    }
  },
  "peace-time-config": {
  "idle-config": {
    "heartbeat-interval": {
      "current-value": 30,
      "min-value": 15,
      "max-value": 240
    },
    "missing-hb-allowed": {
      "current-value": 5,
      "min-value": 3,
      "max-value": 9
    },
    "max-retransmit": {
      "current-value": 3,
      "min-value": 2,
      "max-value": 15
    }
  }
}

```

```

    },
    "ack-timeout": {
      "current-value": 2,
      "min-value": 1,
      "max-value": 30
    },
    "ack-random-factor": {
      "current-value": 1.5,
      "min-value": 1.1,
      "max-value": 4.0
    }
  },
  "trigger-mitigation": {
    "current-value": true
  },
  "config-interval": {
    "current-value": 1439,
    "min-value": 0,
    "max-value": 65535
  }
}
}

```

Figure 17- 18: Configuration response body

#### 4.5.2. Convey DOTS Signal Channel Session Configuration

A PUT request is used to convey the configuration parameters for the signal channel (e.g., heartbeat interval, maximum retransmissions). Message transmission parameters for CoAP are defined in Section 4.8 of [RFC7252]. The RECOMMENDED values of transmission parameter values are ack-timeout (2 seconds), max-retransmit (3), ack-random-factor (1.5). In addition to those parameters, the RECOMMENDED specific DOTS transmission parameter values are 'heartbeat-interval' (30 seconds) and 'missing-hb-allowed' (5).

Note: heartbeat-interval should be tweaked to also assist DOTS messages for NAT traversal (SIG-010 of [I-D.ietf-dots-requirements]). According to [RFC8085], keepalive messages must not be sent more frequently than once every 15 seconds and should use longer intervals when possible. Furthermore, [RFC4787] recommends NATs to use a state timeout of 2 minutes or longer, but experience shows that sending packets every 15 to 30 seconds is necessary to prevent the majority of middleboxes from losing state for UDP flows. From that standpoint, this specification recommends a minimum heartbeat-interval of 15 seconds and a maximum heartbeat-interval of 240 seconds. The recommended value of 30 seconds is selected to anticipate the expiry of NAT state.

A heartbeat-interval of 30 seconds may be seen as too chatty in some deployments. For such deployments, DOTS agents may negotiate longer heartbeat-interval values to prevent any network overload with too frequent keepalives.

When a confirmable "CoAP Ping" is sent, and if there is no response, the "CoAP Ping" is retransmitted max-retransmit number of times by the CoAP layer using an initial timeout set to a random duration between ack-timeout and (ack-timeout\*ack-random-factor) and exponential back-off between retransmissions. By choosing the recommended transmission parameters, the "CoAP Ping" will timeout after 45 seconds. If the DOTS agent does not receive any response from the peer DOTS agent for 'missing-hb-allowed' number of consecutive "CoAP Ping" confirmable messages, it concludes that the DOTS signal channel session is disconnected. A DOTS client MUST NOT transmit a "CoAP Ping" while waiting for the previous "CoAP Ping" response from the same DOTS server.

If the DOTS agent wishes to change the default values of message transmission parameters, then it should follow the guidance given in Section 4.8.1 of [RFC7252]. The DOTS agents MUST use the negotiated values for message transmission parameters and default values for non-negotiated message transmission parameters.

The signal channel session configuration is applicable to a single DOTS signal channel session between the DOTS agents.

```

Header: PUT (Code=0.03)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "config"
Content-Format: "application/cbor"

```

```

{
  "signal-config": {
    "session-id": integer,
"attack-time-config":
    "request-nonce": "string",
    "mitigating-config": {
      "heartbeat-interval": {
        "current-value": integer
      },
      "missing-hb-allowed": {
        "current-value": integer
      },
      "max-retransmit": {
        "current-value": integer
      },
      "ack-timeout": {
        "current-value": integer
      },
      "ack-random-factor": {
        "current-value": number
      }
    },
"peace-time-config":
    "idle-config": {
      "heartbeat-interval": {
        "current-value": integer
      },
      "missing-hb-allowed": {
        "current-value": integer
      },
      "max-retransmit": {
        "current-value": integer
      },
      "ack-timeout": {
        "current-value": integer
      },
      "ack-random-factor": {
        "current-value": number
      }
    },
    "trigger-mitigation": boolean,
    "config-interval": integer
  }
}

```

Figure ~~18~~ 19: PUT to convey the DOTS signal channel session configuration data.

The parameters in Figure ~~18~~ 19 are described below:

**session-id:** Identifier for the DOTS signal channel session configuration data represented as an integer. This identifier MUST be generated by the DOTS client. This document does not make any assumption about how this identifier is generated.

This is a mandatory attribute.

~~attack-time-config:~~

**mitigation-config:** Set of configuration parameters to use when an

~~attack~~ a

**mitigation** is active. The following parameters may be included:

**heartbeat-interval:** Time interval in seconds between two consecutive heartbeat messages.

'0' is used to disable the heartbeat mechanism.

This is an optional attribute.

**missing-hb-allowed:** Maximum number of consecutive heartbeat messages for which the DOTS agent did not receive a response before concluding that the session is disconnected.

This is an optional attribute.

**max-retransmit:** Maximum number of retransmissions for a message (referred to as MAX\_RETRANSMIT parameter in CoAP).

This is an optional attribute.

**ack-timeout:** Timeout value in seconds used to calculate the initial retransmission timeout value (referred to as ACK\_TIMEOUT parameter in CoAP).

This is an optional attribute.

ack-random-factor: Random factor used to influence the timing of retransmissions (referred to as ACK\_RANDOM\_FACTOR parameter in CoAP).

This is an optional attribute.

~~peace-time-config:~~

~~idle-config:~~ Set of configuration parameters to use ~~during~~  
~~peacetime~~, when no  
mitigation is active. This attribute has the same structure as ~~'attack-time-~~  
~~config'~~.  
~~'mitigating-config'~~.

trigger-mitigation: If the parameter value is set to 'false', then DDoS mitigation is triggered only when the DOTS signal channel session is lost. Automated mitigation on loss of signal is discussed in Section 3.3.3 of [I-D.ietf-dots-architecture].

If the DOTS client ceases to respond to heartbeat messages, the DOTS server can detect that the DOTS session is lost.

The default value of the parameter is 'true'.

This is an optional attribute.

config-interval: This parameter is returned to indicate the time interval expressed in minutes, which a DOTS agent must wait for before re-contacting its peer in order to retrieve the signal channel configuration data.

'0' is used to disable this refresh mechanism.

If a non-null value of 'config-interval' is received by a DOTS agent, it has to issue a PUT request to refresh the configuration parameters for the signal channel before the expiry of 'config-interval'. When a DDoS attack is active, refresh requests MUST NOT be sent by DOTS clients and the DOTS server MUST NOT terminate the (D)TLS session after the expiry of 'config-interval'.

This mechanism allows to update the configuration data if a change occurs at the DOTS server side. For example, the new configuration may instruct a DOTS client to cease heartbeats or reduce heartbeat frequency.

If this parameter is not returned, this is equivalent to receiving a 'config-interval' value set to '0'.

If a DOTS server detects that a misbehaving DOTS client does not contact the DOTS server after the expiry of 'config-interval', in order to retrieve the signal channel configuration data, it MAY terminate the (D)TLS session. A (D)TLS session is terminated by the receipt of an authenticated message that closes the connection (e.g., a fatal alert (Section 7.2 of [RFC5246])).

This is an optional attribute.

At least one of the attributes 'heartbeat-interval', 'missing-hb-allowed', 'max-retransmit', 'ack-timeout', 'ack-random-factor', and 'trigger-mitigation' MUST be present in the PUT request. The PUT request with a higher numeric 'session-id' value overrides the DOTS signal channel session configuration data installed by a PUT request with a lower numeric 'session-id' value.

Figure 19 20 shows a PUT request example to convey the configuration parameters for the DOTS signal channel. In this example, heartbeat mechanism is disabled ~~during peacetime~~, when no mitigation is active, while the heartbeat interval is set to '91' when ~~an-attack~~ a mitigation is active.

```
Header: PUT (Code=0.03)
Uri-Host: "www.example.com"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "v1"
Uri-Path: "config"
Content-Format: "application/cbor"
{
  "signal-config": {
    "session-id": 1234534333242,
"attack-time-config":
    "request-nonce": "KseTAjxu6GdFpUROVNhg",
```



```

    "mitigating-config": {
      "heartbeat-interval": {
        "current-value": 91
      },
      "missing-hb-allowed": {
        "current-value": 3
      },
      "max-retransmit": {
        "current-value": 7
      },
      "ack-timeout": {
        "current-value": 5
      },
      "ack-random-factor": {
        "current-value": 1.5
      }
    },
"peace-time-config":
    "idle-config": {
      "heartbeat-interval": {
        "current-value": 0
      },
      "max-retransmit": {
        "current-value": 7
      },
      "ack-timeout": {
        "current-value": 5
      },
      "ack-random-factor": {
        "current-value": 1.5
      }
    },
    "trigger-mitigation": false
  }
}

```

Figure ~~19~~ 20: PUT to convey the configuration parameters

The DOTS server indicates the result of processing the PUT request using CoAP response codes:

- o If the DOTS server finds the 'session-id' parameter value conveyed in the PUT request in its configuration data and if the DOTS server has accepted the updated configuration parameters, then 2.04 (Changed) code is returned in the response.
- o If the DOTS server does not find the 'session-id' parameter value conveyed in the PUT request in its configuration data and if the DOTS server has accepted the configuration parameters, then a response code 2.01 (Created) is returned in the response.
- o If the request is missing one or more mandatory attributes or it contains one or more invalid or unknown parameters, 4.00 (Bad Request) is returned in the response.
- o Response code 4.22 (Unprocessable Entity) is returned in the response, if any of the 'heartbeat-interval', 'missing-hb-allowed', 'max-retransmit', 'target-protocol', 'ack-timeout', and 'ack-random-factor' attribute values are not acceptable to the DOTS server. Upon receipt of the 4.22 error response code, the DOTS client should request the maximum and minimum attribute values acceptable to the DOTS server (Section 4.5.1).

The DOTS client may re-try and send the PUT request with updated attribute values acceptable to the DOTS server.

#### 4.5.3. Delete DOTS Signal Channel Session Configuration

A DELETE request is used to delete the installed DOTS signal channel session configuration data (Figure ~~20~~ 21).

```

Header: DELETE (Code=0.04)
Uri-Host: "host"
Uri-Path: ".well-known"
Uri-Path: "dots"
Uri-Path: "version"
Uri-Path: "config"
Content-Format: "application/cbor"
{
  "signal-config": {
    "session-id": integer,
    "request-nonce": "string"
  }
}

```

Figure ~~20~~ **21**: DELETE configuration

The DOTS server resets the DOTS signal channel session configuration back to the default values and acknowledges a DOTS client's request to remove the DOTS signal channel session configuration using 2.02 (Deleted) response code.

#### 4.6. Redirected Signaling

Redirected DOTS signaling is discussed in detail in Section 3.2.2 of [I-D.ietf-dots-architecture].

If a DOTS server wants to redirect a DOTS client to an alternative DOTS server for a signal session, then the response code 3.00 (alternate server) will be returned in the response to the client.

The DOTS server can return the error response code 3.00 in response to a PUT request from the DOTS client or convey the error response code 3.00 in a unidirectional notification response from the DOTS server.

The DOTS server in the error response conveys the alternate DOTS server's FQDN, and the alternate DOTS server's IP address(es) and time to live values in the CBOR body (Figure ~~21~~ **22**).

```
{
  "alt-server": "string",
  "alt-server-record": [
    {
      "addr": "string",
      "ttl" : integer
    }
  ]
}
```

Figure ~~21~~ **22**: Error response body

The parameters are described below:

alt-server: FQDN of an alternate DOTS server.

addr: IP address of an alternate DOTS server.

ttl: Time to live (TTL) represented as an integer number of seconds.

Figure ~~22~~ **23** shows a 3.00 response example to convey the DOTS alternate server 'alt-server.example', its IP addresses 2001:db8:6401::1 and 2001:db8:6401::2, and TTL values 3600 and 1800.

```
{
  "alt-server": "alt-server.example",
  "alt-server-record": [
    {
      "ttl" : 3600,
      "addr": "2001:db8:6401::1"
    },
    {
      "ttl" : 1800,
      "addr": "2001:db8:6401::2"
    }
  ]
}
```

Figure ~~22~~ **23**: Example of error response body

When the DOTS client receives 3.00 response, it considers the current request as failed, but SHOULD try re-sending the request to the alternate DOTS server. During a DDOS attack, the DNS server may be the target of another DDOS attack, alternate DOTS server's IP addresses conveyed in the 3.00 response help the DOTS client skip DNS lookup of the alternate DOTS server. The DOTS client can then try to establish a UDP or a TCP session with the alternate DOTS server. The DOTS client SHOULD implement a DNS64 function to handle the scenario where an IPv6-only DOTS client communicates with an IPv4-only alternate DOTS server.

#### 4.7. Heartbeat Mechanism

To provide an indication of signal health and distinguish an 'idle' signal channel from a 'disconnected' or 'defunct' session, the DOTS agent sends a heartbeat over the signal channel to maintain its half of the channel. The DOTS agent similarly expects a heartbeat from its peer DOTS agent, and may consider a session terminated in the

prolonged absence of a peer agent heartbeat.

While the communication between the DOTS agents is quiescent, the DOTS client will probe the DOTS server to ensure it has maintained cryptographic state and vice versa. Such probes can also keep firewalls and/or stateful translators bindings alive. This probing reduces the frequency of establishing a new handshake when a DOTS signal needs to be conveyed to the DOTS server.

In order to avoid complications due to the presence of some stateful translators and firewalls (e.g., discard an incoming packet because no matching state is found), DOTS servers MAY trigger their heartbeat requests immediately after receiving heartbeat probes from peer DOTS clients. **It is the responsibility of DOTS clients to ensure that on-path translators are maintaining a binding so that the same external IP address and/or port number are assigned.**

In case of a massive DDoS attack that saturates the incoming link(s) to the DOTS client, all traffic from the DOTS server to the DOTS client will likely be dropped, although the DOTS server receives heartbeat requests in addition to DOTS messages sent by the DOTS client. In this scenario, the DOTS agents MUST behave differently to handle message transmission and DOTS session liveness during link saturation:

- o The DOTS client MUST NOT consider the DOTS session terminated even after a maximum 'missing-hb-allowed' threshold is reached. The DOTS client SHOULD keep on using the current DOTS session to send heartbeat requests over it, so that the DOTS server knows the DOTS client has not disconnected the DOTS session.

After the maximum 'missing-hb-allowed' threshold is reached, the DOTS client SHOULD try to resume the (D)TLS session. The DOTS client SHOULD send mitigation requests over the current DOTS session, and in parallel, for example, try to resume the (D)TLS session or use 0-RTT mode in DTLS 1.3 to piggyback the mitigation request in the ClientHello message.

As soon as the link is no longer saturated, if traffic from the DOTS server reaches the DOTS client over the current DOTS session, the DOTS client can stop (D)TLS session resumption or if (D)TLS session resumption is successful then disconnect the current DOTS session.

- o If the DOTS server does not receive any traffic from the peer DOTS client, then the DOTS server sends heartbeat requests to the DOTS client and after maximum 'missing-hb-allowed' threshold is reached, the DOTS server concludes the session is disconnected.

In DOTS over UDP, heartbeat messages MUST be exchanged between the DOTS agents using the "CoAP Ping" mechanism defined in Section 4.2 of [RFC7252]. Concretely, the DOTS agent sends an Empty Confirmable message and the peer DOTS agent will respond by sending a Reset message.

In DOTS over TCP, heartbeat messages MUST be exchanged between the DOTS agents using the Ping and Pong messages specified in Section 4.4 of [I-D.ietf-core-coap-tcp-tls]. That is, the DOTS agent sends a Ping message and the peer DOTS agent would respond by sending a single Pong message.

## 5. DOTS Signal Channel YANG Module

This document defines a YANG [RFC7950] module for mitigation scope and DOTS signal channel session configuration data.

### 5.1. Tree Structure

This document defines the YANG module "ietf-dots-signal" (Section 5.2), which has the following tree structure. A DOTS signal message can either be a mitigation or a configuration message.

```
module: ietf-dots-signal
  +--rw dots-signal
    +--rw (message-type)?
      +--:(mitigation-scope)
        | +--rw client-identifier* client-domain-hash*  binary
        | +--rw request-nonce* string
        | +--rw scope* [mitigation-id]
        |   +--rw mitigation-id int32
        |   +--rw target-prefix* inet:ip-prefix
        |   +--rw target-port-range* [lower-port upper-port]
        |     | +--rw lower-port inet:port-number
        |     | +--rw upper-port inet:port-number
```

```

|--rw target-protocol*      uint8
|--rw target-fqdn*          inet:domain-name
|--rw target-uri*           inet:uri
|--rw alias-name*           string
|--rw lifetime?             int32
|--rw mitigation-start?     int64
|--rw status?               enumeration
+--ro conflict-information
| +--ro conflict-status?    enumeration
| +--ro conflict-cause?    enumeration
| +--ro retry-timer?       int32
| +--ro conflict-scope
| | +--ro target-prefix*    inet:ip-prefix
| | +--ro target-port-range* [lower-port upper-port]
| | | +--ro lower-port      inet:port-number
| | | +--ro upper-port      inet:port-number
| | +--ro target-protocol*  uint8
| | +--ro target-fqdn*      inet:domain-name
| | +--ro target-uri*       inet:uri
| | +--ro alias-name*       string
| | +--ro acl-list* [acl-name acl-type]
| | | +--ro acl-name        -> /ietf-acl:access-lists/acl/acl-name
| | | +--ro acl-type        -> /ietf-acl:access-lists/acl/acl-type
+--ro pkts-dropped?         yang:zero-based-counter64
+--ro bps-dropped?          yang:zero-based-counter64
+--ro bytes-dropped?        yang:zero-based-counter64
+--ro pps-dropped?          yang:zero-based-counter64
+---:(configuration)
+---:(signal-config)
|--rw session-id            int32
+--rw attack-time-config request-nonce      string
+--rw mitigating-config
| +--rw heartbeat-interval
| | +--rw max-value?      int16
| | +--rw min-value?      int16
| | +--rw current-value?  int16
+--rw missing-hb-allowed
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw max-retransmit
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw ack-timeout
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw ack-random-factor
| +--rw max-value?      decimal64
| +--rw min-value?      decimal64
| +--rw current-value?  decimal64
+--rw peace-time-config idle-config
| +--rw heartbeat-interval
| | +--rw max-value?      int16
| | +--rw min-value?      int16
| | +--rw current-value?  int16
+--rw missing-hb-allowed
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw max-retransmit
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw ack-timeout
| +--rw max-value?      int16
| +--rw min-value?      int16
| +--rw current-value?  int16
+--rw ack-random-factor
| +--rw max-value?      decimal64
| +--rw min-value?      decimal64
| +--rw current-value?  decimal64
+--rw trigger-mitigation?  boolean
+--rw config-interval?     int32

```

## 5.2. YANG Module

<CODE BEGINS> file "ietf-dots-signal@2017-12-19.yang"

```

module ietf-dots-signal {
  yang-version 1.1;
  namespace "urn:ietf:params:xml:ns:yang:ietf-dots-signal";
  prefix "signal";

```

```
import ietf-inet-types {prefix "inet";}
import ietf-yang-types {prefix yang;}
import ietf-access-control-list {prefix "ietf-acl";}

organization "IETF DDoS Open Threat Signaling (DOTS) Working Group";

contact
  "Konda, Tirumaleswar Reddy <TirumaleswarReddy_Konda@McAfee.com>
  Mohamed Boucadair <mohamed.boucadair@orange.com>
  Prashanth Patil <praspati@cisco.com>
  Andrew Mortensen <amortensen@arbor.net>
  Nik Teague <nteague@verisign.com>";

description
  "This module contains YANG definition for the signaling
  messages exchanged between a DOTS client and a DOTS server.

  Copyright (c) 2017 IETF Trust and the persons identified as
  authors of the code. All rights reserved.

  Redistribution and use in source and binary forms, with or
  without modification, is permitted pursuant to, and subject
  to the license terms contained in, the Simplified BSD License
  set forth in Section 4.c of the IETF Trust's Legal Provisions
  Relating to IETF Documents
  (http://trustee.ietf.org/license-info).

  This version of this YANG module is part of RFC XXXX; see
  the RFC itself for full legal notices.";

revision 2017-12-19 {
  description
    "Initial revision.";
  reference
    "RFC XXXX: Distributed Denial-of-Service Open Threat
    Signaling (DOTS) Signal Channel";
}

grouping target {
  description
    "Specifies the scope of the mitigation request.";

  leaf-list target-prefix {
    type inet:ip-prefix;
    description
      "IPv4 or IPv6 prefix identifying the target.";
  }

  list target-port-range {
    key "lower-port upper-port";

    description
      "Port range. When only lower-port is
      present, it represents a single port.";

    leaf lower-port {
      type inet:port-number;
      mandatory true;
      description "Lower port number.";
    }

    leaf upper-port {
      type inet:port-number;
      must ". >= ../lower-port" {
        error-message
          "The upper port number must be greater than
          or equal to lower port number.";
      }
      description "Upper port number.";
    }
  }

  leaf-list target-protocol {
    type uint8;
    description
      "Identifies the target protocol number.

      The value '0' means 'all protocols'.

      Values are taken from the IANA protocol registry:
      https://www.iana.org/assignments/protocol-numbers/protocol-numbers.xhtml

      For example, 6 for TCP or 17 for UDP.";
```

```

}
leaf-list target-fqdn {
  type inet:domain-name;
  description "FQDN identifying the target.";
}

leaf-list target-uri {
  type inet:uri;
  description "URI identifying the target.";
}

leaf-list alias-name {
  type string;
  description "alias name";
}
}

grouping mitigation-scope {
  description
    "Specifies the scope of the mitigation request.";

  leaf-list client-identifier client-domain-hash {
    type binary;
    description
      "The client identifier domain hash may be conveyed by
      the server-side DOTS gateway to propagate the DOTS
      client domain identification information from the
      gateway's client-side to the gateway's server-side,
      and from the gateway's server-side to the DOTS
      server.

      It may be used by the final DOTS server
      for policy enforcement purposes.";
  }

  leaf request-nonce {
    type string;
    mandatory true;
    description
      "A globally unique identifier that is randomly
      generated by a DOTS client to prevent
      request collisions.";
  }

  list scope {
    key mitigation-id;
    description
      "The scope of the request.";

    leaf mitigation-id {
      type int32;
      description
        "Mitigation request identifier.

        This identifier must be unique for each mitigation
        request bound to the DOTS client.";
    }
  }

  uses target;

  leaf lifetime {
    type int32;
    units "seconds";
    default 3600;
    description
      "Indicates the lifetime of the mitigation request.";
    reference
      "RFC XXXX: Distributed Denial-of-Service Open Threat
      Signaling (DOTS) Signal Channel";
  }

  leaf mitigation-start {
    type int64;
    units "seconds";
    description
      "Mitigation start time is represented in seconds
      relative to 1970-01-01T00:00Z in UTC time.";
  }

  leaf status {
    type enumeration {
      enum "attack-mitigation-in-progress" {
        value 1;
        description

```

```
        "Attack mitigation is in progress (e.g., changing
        the network path to re-route the inbound traffic
        to DOTS mitigator).";
    }

    enum "attack-successfully-mitigated" {
        value 2;
        description
            "Attack is successfully mitigated (e.g., traffic
            is redirected to a DDOS mitigator and attack
            traffic is dropped or blackholed).";
    }
    enum "attack-stopped" {
        value 3;
        description
            "Attack has stopped and the DOTS client can
            withdraw the mitigation request.";
    }

    enum "attack-exceeded-capability" {
        value 4;
        description
            "Attack has exceeded the mitigation provider
            capability.";
    }

    enum "dots-client-withdrawn-mitigation" {
        value 5;
        description
            "DOTS client has withdrawn the mitigation
            request and the mitigation is active but
            terminating.";
    }

    enum "attack-mitigation-terminated" {
        value 6;
        description
            "Attack mitigation is now terminated.";
    }

    enum "attack-mitigation-withdrawn" {
        value 7;
        description
            "Attack mitigation is withdrawn.";
    }

    enum "attack-mitigation-rejected" {
        value 8;
        description
            "Attack mitigation is rejected.";
    }
}
config false;
description
    "Indicates the status of a mitigation request.
    It must be included in responses only.";
}

container conflict-information {
    config false;
    description
        "Indicates that a conflict is detected.
        Must only be used for responses.";
}

leaf conflict-status {
    type enumeration {
        enum "request-inactive-other-active" {
            value 1;
            description
                "DOTS Server has detected conflicting mitigation
                requests from different DOTS clients.
                This mitigation request is currently inactive
                until the conflicts are resolved. Another
                mitigation request is active.";
        }

        enum "request-active" {
            value 2;
            description
                "DOTS Server has detected conflicting mitigation
                requests from different DOTS clients.
                This mitigation request is currently active.";
        }
    }
}
```

```

        enum "all-requests-inactive" {
            value 3;
            description
                "DOTS Server has detected conflicting mitigation
                 requests from different DOTS clients. All
                 conflicting mitigation requests are inactive.";
        }
    }
    description
        "Indicates the conflict status.
         It must be included in responses only.";
}

leaf conflict-cause {
    type enumeration {
        enum "overlapping-targets" {
            value 1;
            description
                "Overlapping targets. conflict-scope provides
                 more details about the exact conflict.";
        }

        enum "conflict-with-whitelist" {
            value 2;
            description
                "Conflicts with an existing white list.

                 This code is returned when the DDoS mitigation
                 detects that some of the source addresses/prefixes
                 listed in the white list ACLs are actually
                 attacking the target.";
        }
    }
    description
        "Indicates the cause of the conflict.
         It must be included in responses only.";
}

leaf retry-timer {
    type int32;
    units "seconds";
    description
        "The DOTS client must not re-send the
         same request before the expiry of this timer.
         It must be included in responses, only.";
}

container conflict-scope {
    description
        "Provides more information about the conflict scope.";

    uses target {
        when "../conflict-cause = 'overlapping-targets'";
    }

    list acl-list {
        when "../conflict-cause = 'conflict-with-whitelist'";
        key "acl-name acl-type";
        description
            "List of conflicting ACLs";

        leaf acl-name {
            type leafref {
                path "/ietf-acl:access-lists/ietf-acl:acl" +
                    "/ietf-acl:acl-name";
            }
            description
                "Reference to the conflicting ACL name bound to
                 a DOTS client.";
        }

        leaf acl-type {
            type leafref {
                path "/ietf-acl:access-lists/ietf-acl:acl" +
                    "/ietf-acl:acl-type";
            }
            description
                "Reference to the conflicting ACL type bound to
                 a DOTS client.";
        }
    }
}
}
}

```



```

    leaf pkts-dropped {
      type yang:zero-based-counter64;
      config false;
      description
        "Number of dropped packets";
    }

    leaf bps-dropped {
      type yang:zero-based-counter64;
      config false;
      description
        "The average number of dropped bytes per second for
        the mitigation request since the attack
        mitigation is triggered.";
    }

    leaf bytes-dropped {
      type yang:zero-based-counter64;
      units 'bytes';
      config false;
      description
        "Counter for dropped packets; in bytes.";
    }

    leaf pps-dropped {
      type yang:zero-based-counter64;
      config false;
      description
        "The average number of dropped packets per second
        for the mitigation request since the attack
        mitigation is triggered.";
    }
  }
}

grouping config-parameters {
  description
    "Subset of DOTS signal channel session configuration.";

  container heartbeat-interval {
    description
      "DOTS agents regularly send heartbeats to each other
      after mutual authentication is successfully
      completed in order to keep the DOTS signal channel
      open.";

    leaf max-value {
      type int16;
      units "seconds";
      description
        "Maximum acceptable value.";
      reference
        "RFC XXXX: Distributed Denial-of-Service Open Threat
        Signaling (DOTS) Signal Channel";
    }

    leaf min-value {
      type int16;
      units "seconds";
      description
        "Minimum acceptable value.";
      reference
        "RFC XXXX: Distributed Denial-of-Service Open Threat
        Signaling (DOTS) Signal Channel";
    }
  }

  leaf current-value {
    type int16;
    units "seconds";
    default 30;
    description
      "Current value.

      '0' means that heartbeat mechanism is deactivated.";
    reference
      "RFC XXXX: Distributed Denial-of-Service Open Threat
      Signaling (DOTS) Signal Channel";
  }
}

container missing-hb-allowed {
  description
    "Maximum number of missing heartbeats allowed.";

  leaf max-value {

```

```
    type int16;
    description
        "Maximum acceptable value.";
    reference
        "RFC XXXX: Distributed Denial-of-Service Open Threat
        Signaling (DOTS) Signal Channel";
}

leaf min-value {
    type int16;
    description
        "Minimum acceptable value.";
    reference
        "RFC XXXX: Distributed Denial-of-Service Open Threat
        Signaling (DOTS) Signal Channel";
}

leaf current-value {
    type int16;
    default 5;
    description
        "Current value.";
    reference
        "RFC XXXX: Distributed Denial-of-Service Open Threat
        Signaling (DOTS) Signal Channel";
}
}

container max-retransmit {
    description
        "Maximum number of retransmissions of a Confirmable
        message.";

    leaf max-value {
        type int16;
        description
            "Maximum acceptable value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf min-value {
        type int16;
        description
            "Minimum acceptable value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf current-value {
        type int16;
        default 3;
        description
            "Current value.";
        reference
            "RFC XXXX: Distributed Denial-of-Service Open Threat
            Signaling (DOTS) Signal Channel";
    }
}

container ack-timeout {
    description
        "Initial retransmission timeout value.";

    leaf max-value {
        type int16;
        units "seconds";
        description
            "Maximum value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf min-value {
        type int16;
        units "seconds";
        description
            "Minimum value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf current-value {
        type int16;
        units "seconds";
        default 2;
        description
```

```

        "Current value.";
    reference
        "Section 4.8 of RFC 7552.";
}
}

container ack-random-factor {
    description
        "Random factor used to influence the timing of
        retransmissions.";

    leaf max-value {
        type decimal64 {
            fraction-digits 2;
        }
        description
            "Maximum acceptable value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf min-value {
        type decimal64 {
            fraction-digits 2;
        }
        description
            "Minimum acceptable value.";
        reference
            "Section 4.8 of RFC 7552.";
    }

    leaf current-value {
        type decimal64 {
            fraction-digits 2;
        }
        default 1.5;
        description
            "Current value.";
        reference
            "Section 4.8 of RFC 7552.";
    }
}

}

grouping signal-config {
    description
        "DOTS signal channel session configuration.";

    leaf session-id {
        type int32;
        mandatory true;
        description
            "An identifier for the DOTS signal channel
            session configuration data.";
    }

    leaf request-nonce {
        type string;
        mandatory true;
        description
            "A globally unique identifier that is meant to
            avoid collision.";
    }

    container attack-time-config mitigating-config {
        description
            "Configuration paramaters parameters to use when an-attack a mitigation is active.";
        uses config-parameters;
    }

    container peace-time-config idle-config {
        description
            "Configuration paramaters parameters to use in-peace-time. when no mitigation is
            active.";
        uses config-parameters;
    }

    leaf trigger-mitigation {
        type boolean;
        default true;
        description
            "If false, then mitigation is triggered
            only when the DOTS server channel session is lost";
        reference
            "RFC XXXX: Distributed Denial-of-Service Open Threat

```

```
        Signaling (DOTS) Signal Channel";
    }

    leaf config-interval {
        type int32;
        units "minutes";
        description
            "This parameter is returned by a DOTS server to
            a requesting DOTS client to indicate the time interval
            after which the DOTS client must contact the DOTS
            server in order to retrieve the signal channel
            configuration data.

            This mechanism allows the update of the configuration
            data if a change occurs.

            For example, the new configuration may instruct
            a DOTS client to cease heartbeats or reduce
            heartbeat frequency.

            '0' is used to disable this refresh mechanism.";
    }
}
container dots-signal {
    description
        "Main container for DOTS signal message.
        A DOTS signal message can be a mitigation message or
        a configuration message.";

    choice message-type {
        description
            "Either a mitigation or a configuration message.";

        case mitigation-scope {
            description
                "Mitigation scope of a mitigation message.";
            uses mitigation-scope;
        }

        case configuration signal-config {
            description
                "Configuration message.";
            uses signal-config;
        }
    }
}
}
<CODE ENDS>
```

## 6. Mapping Parameters to CBOR

All parameters in the payload of the DOTS signal channel MUST be mapped to CBOR types as shown in Table 4 and are assigned an integer key to save space. The recipient of the payload MAY reject the information if it is not suitably mapped.

Parameter name	CBOR key	CBOR major type of value
mitigation-scope	1	5 (map)
scope	2	5 (map)
mitigation-id	3	0 (unsigned)
acl-list	4	4
target-port-range	5	4
lower-port	6	0
upper-port	7	0
target-protocol	8	4
target-fqdn	9	4
target-uri	10	4
alias-name	11	4
lifetime	12	0
attack-status	13	0
signal-config	14	5
heartbeat-interval	15	5 (map)
max-retransmit	16	5 (map)
ack-timeout	17	5 (map)
ack-random-factor	18	5 (map)
min-value	19	0
max-value	20	0
status	21	0
conflict-information	22	5 (map)
conflict-status	23	0
conflict-cause	24	0
retry-timer	25	0

bytes-dropped	26	0		
bps-dropped	27	0		
pkts-dropped	28	0		
pps-dropped	29	0		
session-id	30	0		
trigger-mitigation	31	7 (simple types)		
missing-hb-allowed	32	5 (map)		
current-value	33	0		
mitigation-start	34	7 (floating-point)		
target-prefix	35	4 (array)		
<del>client-identifier</del> <b>client-domain-hash</b>		36	2 (byte string)	
alt-server	37	2		
alt-server-record	38	4		
addr	39	2		
ttl	40	0		
conflict-scope	41	5 (map)		
acl-name	42	3		
acl-type	43	3		
config-interval	44	0		
<del>attack-time-config</del> <b>mitigating-config</b>		45	5 (map)	
<del>peace-time-config</del> <b>idle-config</b>		46	5 (map)	
<b>request-nonce</b>	47	2 (byte string)		

Table 4: CBOR mappings used in DOTS signal channel message

## 7. (D)TLS Protocol Profile and Performance Considerations

### 7.1. (D)TLS Protocol Profile

This section defines the (D)TLS protocol profile of DOTS signal channel over (D)TLS and DOTS data channel over TLS.

There are known attacks on (D)TLS, such as man-in-the-middle and protocol downgrade attacks. These are general attacks on (D)TLS and, as such, they are not specific to DOTS over (D)TLS; refer to the (D)TLS RFCs for discussion of these security issues. DOTS agents MUST adhere to the (D)TLS implementation recommendations and security considerations of [RFC7525] except with respect to (D)TLS version. Since DOTS signal channel encryption relies upon (D)TLS is virtually a green-field deployment, DOTS agents MUST implement only (D)TLS 1.2 or later.

When a DOTS client is configured with a domain name of the DOTS server, and connects to its configured DOTS server, the server may present it with a PKIX certificate. In order to ensure proper authentication, a DOTS client MUST verify the entire certification path per [RFC5280]. The DOTS client additionally uses [RFC6125] validation techniques to compare the domain name with the certificate provided.

A key challenge to deploying DOTS is the provisioning of DOTS clients, including the distribution of keying material to DOTS clients to enable the required mutual authentication of DOTS agents. EST defines a method of certificate enrollment by which domains operating DOTS servers may provide DOTS clients with all the necessary cryptographic keying material, including a private key and a certificate to authenticate themselves. One deployment option is DOTS clients behave as EST clients for certificate enrollment from an EST server provisioned by the mitigation provider. This document does not specify which EST mechanism the DOTS client uses to achieve initial enrollment.

Implementations compliant with this profile MUST implement all of the following items:

- o DTLS record replay detection (Section 3.3 of [RFC6347]) to protect against replay attacks.
- o (D)TLS session resumption without server-side state [RFC5077] to resume session and convey the DOTS signal.
- o Raw public keys [RFC7250] or PSK handshake [RFC4279] which reduces the size of the ServerHello, and can be used by DOTS agents that cannot obtain certificates.

Implementations compliant with this profile SHOULD implement all of the following items to reduce the delay required to deliver a DOTS signal channel message:

- o TLS False Start [RFC7918] which reduces round-trips by allowing the TLS second flight of messages (ChangeCipherSpec) to also contain the DOTS signal.
- o Cached Information Extension [RFC7924] which avoids transmitting

the server's certificate and certificate chain if the client has cached that information from a previous TLS handshake.

- o TCP Fast Open [RFC7413] can reduce the number of round-trips to convey DOTS signal channel message.

## 7.2. (D)TLS 1.3 Considerations

TLS 1.3 [I-D.ietf-tls-tls13] provides critical latency improvements for connection establishment over TLS 1.2. The DTLS 1.3 protocol [I-D.ietf-tls-dtls13] is based upon the TLS 1.3 protocol and provides equivalent security guarantees. (D)TLS 1.3 provides two basic handshake modes the DOTS signal channel can take advantage of:

- o A full handshake mode in which a DOTS client can send a DOTS mitigation request message after one round trip and the DOTS server immediately responds with a DOTS mitigation response. This assumes no packet loss is experienced.
- o 0-RTT mode in which the DOTS client can authenticate itself and send DOTS mitigation request messages in the first message, thus reducing handshake latency. 0-RTT only works if the DOTS client has previously communicated with that DOTS server, which is very likely with the DOTS signal channel.

The DOTS client has to establish a (D)TLS session with the DOTS server during peacetime and share a PSK.

During a DDoS attack, the DOTS client can use the (D)TLS session to convey the DOTS mitigation request message and, if there is no response from the server after multiple retries, the DOTS client can resume the (D)TLS session in 0-RTT mode using PSK.

Section 8 of [I-D.ietf-tls-tls13] discusses some mechanisms to implement to limit the impact of replay attacks on 0-RTT data. If TLS1.3 is used, DOTS servers must implement one of these mechanisms.

A simplified TLS 1.3 handshake with 0-RTT DOTS mitigation request message exchange is shown in Figure 23- 24.



Figure 23- 24: TLS 1.3 handshake with 0-RTT

## 7.3. MTU and Fragmentation

To avoid DOTS signal message fragmentation and the subsequent decreased probability of message delivery, DOTS agents MUST ensure that the DTLS record MUST fit within a single datagram. If the path MTU is not known to the DOTS server, an IP MTU of 1280 bytes SHOULD be assumed. The length of the URL MUST NOT exceed 256 bytes. If UDP is used to convey the DOTS signal messages then the DOTS client must consider the amount of record expansion expected by the DTLS processing when calculating the size of CoAP message that fits within the path MTU. Path MTU MUST be greater than or equal to [CoAP message size + DTLS overhead of 13 octets + authentication overhead of the negotiated DTLS cipher suite + block padding (Section 4.1.1.1 of [RFC6347])]. If the request size exceeds the path MTU then the DOTS client MUST split the DOTS signal into separate messages, for example the list of addresses in the 'target-prefix' parameter could be split into multiple lists and each list conveyed in a new PUT request.

Implementation Note: DOTS choice of message size parameters works well with IPv6 and with most of today's IPv4 paths. However, with IPv4, it is harder to reliably ensure that there is no IP fragmentation. If IPv4 path MTU is unknown, implementations may want to limit themselves to more conservative IPv4 datagram sizes such as

576 bytes, as per [RFC0791]. IP packets whose size does not exceed 576 bytes should never need to be fragmented: therefore, sending a maximum of 500 bytes of DOTS signal over a UDP datagram will generally avoid IP fragmentation.

## 8. Mutual Authentication of DOTS Agents & Authorization of DOTS Clients

(D)TLS based upon client certificate can be used for mutual authentication between DOTS agents. If a DOTS gateway is involved, DOTS clients and DOTS gateways MUST perform mutual authentication; only authorized DOTS clients are allowed to send DOTS signals to a DOTS gateway. The DOTS gateway and the DOTS server MUST perform mutual authentication; a DOTS server only allows DOTS signal channel messages from an authorized DOTS gateway, thereby creating a two-link chain of transitive authentication between the DOTS client and the DOTS server.

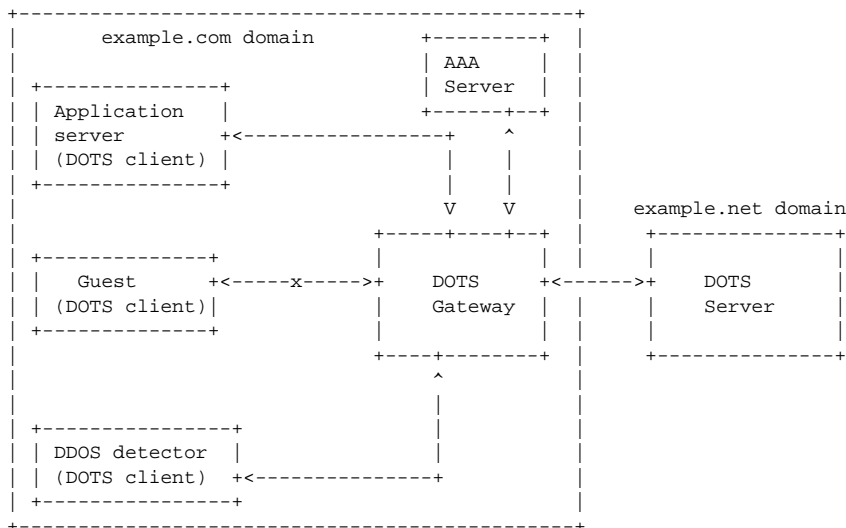


Figure 24- 25: Example of Authentication and Authorization of DOTS Agents

In the example depicted in Figure 24, 25, the DOTS gateway and DOTS clients within the 'example.com' domain mutually authenticate with each other. After the DOTS gateway validates the identity of a DOTS client, it communicates with the AAA server in the 'example.com' domain to determine if the DOTS client is authorized to request DDoS mitigation. If the DOTS client is not authorized, a 4.01 (Unauthorized) is returned in the response to the DOTS client. In this example, the DOTS gateway only allows the application server and DDoS attack detector to request DDoS mitigation, but does not permit the user of type 'guest' to request DDoS mitigation.

Also, DOTS gateways and servers located in different domains MUST perform mutual authentication (e.g., using certificates). A DOTS server will only allow a DOTS gateway with a certificate for a particular domain to request mitigation for that domain. In reference to Figure 24, 25, the DOTS server only allows the DOTS gateway to request mitigation for 'example.com' domain and not for other domains.

## 9. IANA Considerations

This specification registers a service port (Section 9.1), an URI suffix in the Well-Known URIs registry (Section 9.2), a CoAP response code (Section 9.3), a YANG module (Section 9.5). It also creates a registry for mappings to CBOR (Section 9.4).

### 9.1. DOTS Signal Channel UDP and TCP Port Number

IANA is requested to assign the port number TBD to the DOTS signal channel protocol for both UDP and TCP from the "Service Name and Transport Protocol Port Number Registry" available at <https://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>.

The assignment of port number 4646 is strongly suggested, as 4646 is the ASCII decimal value for ".." (DOTS).

### 9.2. Well-Known 'dots' URI

This document requests IANA to register the 'dots' well-known URI in the Well-Known URIs registry (<https://www.iana.org/assignments/well-known-uris/well-known-uris.xhtml>) as defined by [RFC5785].

URI suffix: dots

Change controller: IETF

Specification document(s): This RFC

Related information: None

### 9.3. CoAP Response Code

IANA is requested to add the following entry to the "CoAP Response Codes" sub-registry available at <https://www.iana.org/assignments/core-parameters/core-parameters.xhtml#response-codes>:

Code	Description	Reference
3.00	Alternate server	[RFCXXXX]

Table 4: CoAP Response Code

### 9.4. DOTS Signal Channel CBOR Mappings Registry

The document requests IANA to create a new registry, entitled "DOTS Signal Channel CBOR Mappings Registry". The structure of this registry is provided in Section 9.4.1.

The registry is initially populated with the values in Section 9.4.2.

Values from that registry MUST be assigned via Expert Review [RFC8126].

#### 9.4.1. Registration Template

Parameter name:

Parameter name as used in the DOTS signal channel.

CBOR Key Value:

Key value for the parameter. The key value MUST be an integer in the 1-65536 range. The key values in the 32758-65536 range are assigned to Vendor-Specific parameters.

CBOR Major Type:

CBOR Major type and optional tag for the claim.

Change Controller:

For Standards Track RFCs, list the "IESG". For others, give the name of the responsible party. Other details (e.g., postal address, email address, home page URI) may also be included.

Specification Document(s):

Reference to the document or documents that specify the parameter, preferably including URIs that can be used to retrieve copies of the documents. An indication of the relevant sections may also be included but is not required.

#### 9.4.2. Initial Registry Contents

- o Parameter Name: mitigation-scope
- o CBOR Key Value: 1
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: scope
- o CBOR Key Value: 2
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: mitigation-id
- o CBOR Key Value: 3
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: acl-list
- o CBOR Key Value: 4
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document



- o Parameter Name: target-port-range
- o CBOR Key Value: 5
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: lower-port
- o CBOR Key Value: 6
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: upper-port
- o CBOR Key Value: 7
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: target-protocol
- o CBOR Key Value: 8
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: target-fqdn
- o CBOR Key Value: 9
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: target-uri
- o CBOR Key Value: 10
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: alias-name
- o CBOR Key Value: 11
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: lifetime
- o CBOR Key Value: 12
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: attack-status
- o CBOR Key Value: 13
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: signal-config
- o CBOR Key Value: 14
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: heartbeat-interval
- o CBOR Key Value: 15
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: max-retransmit
- o CBOR Key Value: 16
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ack-timeout
- o CBOR Key Value: 17
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ack-random-factor
- o CBOR Key Value: 18
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: min-value
- o CBOR Key Value: 19
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: max-value
- o CBOR Key Value: 20
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: status
- o CBOR Key Value: 21
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: conflict-information
- o CBOR Key Value: 22
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: conflict-status
- o CBOR Key Value: 23
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: conflict-cause
- o CBOR Key Value: 24
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: retry-timer
- o CBOR Key Value: 25
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: bytes-dropped
- o CBOR Key Value: 26
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: bps-dropped
- o CBOR Key Value: 27
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: pkts-dropped
- o CBOR Key Value: 28
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: pps-dropped
- o CBOR Key Value: 29
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: session-id
- o CBOR Key Value: 30
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: trigger-mitigation
- o CBOR Key Value: 31
- o CBOR Major Type: 7
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: missing-hb-allowed
- o CBOR Key Value: 32
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: current-value
- o CBOR Key Value: 33
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: mitigation-start
- o CBOR Key Value: 34
- o CBOR Major Type: 7
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: target-prefix
- o CBOR Key Value: 35
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ~~client-identifier~~ client-domain-hash
- o CBOR Key Value: 36
- o CBOR Major Type: 2
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: alt-server
- o CBOR Key Value: 37
- o CBOR Major Type: 2
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: alt-server-record
- o CBOR Key Value: 38
- o CBOR Major Type: 4
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: addr
- o CBOR Key Value: 39
- o CBOR Major Type: 2
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ttl
- o CBOR Key Value: 40
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: conflict-scope
- o CBOR Key Value: 41
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: acl-name
- o CBOR Key Value: 42
- o CBOR Major Type: 3
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: acl-type
- o CBOR Key Value: 43
- o CBOR Major Type: 3
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: config-interval
- o CBOR Key Value: 44
- o CBOR Major Type: 0
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ~~attack-time-config~~ mitigating-config
- o CBOR Key Value: 45
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document
  
- o Parameter Name: ~~peace-time-config~~ idle-config
- o CBOR Key Value: 46
- o CBOR Major Type: 5
- o Change Controller: IESG
- o Specification Document(s): this document

- o Parameter Name: request-nonce
- o CBOR Key Value: 47
- o CBOR Major Type: 2
- o Change Controller: IESG
- o Specification Document(s): this document

## 9.5. DOTS Signal Channel YANG Module

This document requests IANA to register the following URI in the "IETF XML Registry" [RFC3688]:

URI: urn:ietf:params:xml:ns:yang:ietf-dots-signal  
Registrant Contact: The IESG.  
XML: N/A; the requested URI is an XML namespace.

This document requests IANA to register the following YANG module in the "YANG Module Names" registry [RFC7950].

name: ietf-signal  
namespace: urn:ietf:params:xml:ns:yang:ietf-dots-signal  
prefix: signal  
reference: RFC XXXX

## 10. Implementation Status

[Note to RFC Editor: Please remove this section and reference to [RFC7942] prior to publication.]

This section records the status of known implementations of the protocol defined by this specification at the time of posting this Internet-Draft, and is based upon a proposal described in [RFC7942]. The description of implementations in this section is intended to assist the IETF in its decision-making process when progressing drafts to RFCs. Please note that the listing of any individual implementation here does not imply endorsement by the IETF. Furthermore, no effort has been spent to verify the information presented here, and which was provided by individuals. This is not intended as, and must not be construed to be, a catalog of available implementations or features. Readers are advised to note that other implementations may exist.

According to [RFC7942], "this will allow reviewers and working groups to assign due consideration to documents that have the benefit of running code, which may serve as evidence of valuable experimentation and feedback that have made the implemented protocols more mature. It is up to the individual working groups to use this information as they see fit".

### 10.1. nttdots

Organization: NTT Communication is developing a DOTS client and DOTS server software based on DOTS signal channel specified in this draft. It will be open-sourced.

Description: Early implementation of DOTS protocol. It is aimed to implement a full DOTS protocol specification in accordance with the nurturing DOTS protocol.

Implementation: <https://github.com/nttdots/go-dots>

Level of maturity: It is an early implementation of the DOTS protocol. Messaging between DOTS clients and DOTS servers has been tested. Level of maturity will increase in accordance with the nurturing DOTS protocol.

Coverage: Capability of DOTS client: sending DOTS messages to the DOTS server in CoAP over DTLS as dots-signal. Capability of DOTS server: receiving dots-signal, validating received dots-signal, starting mitigation by handing over the dots-signal to DDOS mitigator.

Licensing: It will be open-sourced with BSD 3-clause license.

Implementation experience: It is implemented in Go-lang. Core specification of signaling is mature to be implemented, however, finding good libraries (like DTLS, CoAP) is rather difficult.

Contact: Kaname Nishizuka <[kaname@nttv6.jp](mailto:kaname@nttv6.jp)>

## 11. Security Considerations

Authenticated encryption MUST be used for data confidentiality and message integrity. The interaction between the DOTS agents requires Datagram Transport Layer Security (DTLS) and Transport Layer Security (TLS) with a cipher suite offering confidentiality protection and the guidance given in [RFC7525] MUST be followed to avoid attacks on (D)TLS. The (D)TLS protocol profile for DOTS signal channel is specified in Section 7.

A single DOTS signal channel between DOTS agents can be used to exchange multiple DOTS signal messages. To reduce DOTS client and

DOTS server workload, DOTS clients SHOULD re-use the (D)TLS session.

If TCP is used between DOTS agents, an attacker may be able to inject RST packets, bogus application segments, etc., regardless of whether TLS authentication is used. Because the application data is TLS protected, this will not result in the application receiving bogus data, but it will constitute a DoS on the connection. This attack can be countered by using TCP-AO [RFC5925]. If TCP-AO is used, then any bogus packets injected by an attacker will be rejected by the TCP-AO integrity check and therefore will never reach the TLS layer.

In order to prevent leaking internal information outside a client-domain, DOTS gateways located in the client-domain SHOULD NOT reveal the identification information that pertains to internal DOTS clients (~~client-identifier~~) (e.g., source IP address, client's hostname) unless explicitly configured to do so.

Special care should be taken in order to ensure that the activation of the proposed mechanism will not impact the stability of the network (including connectivity and services delivered over that network).

## 12. Contributors

The following individuals have contributed to this document:

Mike Geller Cisco Systems, Inc. 3250 Florida 33309 USA Email: mgeller@cisco.com

Robert Moskowitz HTT Consulting Oak Park, MI 42837 United States Email: rgm@htt-consult.com

Dan Wing Email: dwing-ietf@fuggles.com

## 13. Acknowledgements

Thanks to Christian Jacquenet, Roland Dobbins, Roman D. Danyliw, Michael Richardson, Ehud Doron, Kaname Nishizuka, Dave Dolson, Liang Xia, Gilbert Clark, and Nesredien Suleiman for the discussion and comments.

Special thanks to Jon Shallow for the careful reviews and inputs that enhanced this specification.

## 14. References

### 14.1. Normative References

- [I-D.ietf-core-coap-tcp-tls]  
Bormann, C., Lemay, S., Tschofenig, H., Hartke, K., Silverajan, B., and B. Raymor, "CoAP (Constrained Application Protocol) over TCP, TLS, and WebSockets", draft-ietf-core-coap-tcp-tls-11 (work in progress), December 2017.
- [RFC2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC3688] Mealling, M., "The IETF XML Registry", BCP 81, RFC 3688, DOI 10.17487/RFC3688, January 2004, <<https://www.rfc-editor.org/info/rfc3688>>.
- [RFC4279] Eronen, P., Ed. and H. Tschofenig, Ed., "Pre-Shared Key Ciphersuites for Transport Layer Security (TLS)", RFC 4279, DOI 10.17487/RFC4279, December 2005, <<https://www.rfc-editor.org/info/rfc4279>>.
- [RFC5246] Dierks, T. and E. Rescorla, "The Transport Layer Security (TLS) Protocol Version 1.2", RFC 5246, DOI 10.17487/RFC5246, August 2008, <<https://www.rfc-editor.org/info/rfc5246>>.
- [RFC5280] Cooper, D., Santesson, S., Farrell, S., Boeyen, S., Housley, R., and W. Polk, "Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile", RFC 5280, DOI 10.17487/RFC5280, May 2008, <<https://www.rfc-editor.org/info/rfc5280>>.
- [RFC5785] Nottingham, M. and E. Hammer-Lahav, "Defining Well-Known Uniform Resource Identifiers (URIs)", RFC 5785, DOI 10.17487/RFC5785, April 2010,

<<https://www.rfc-editor.org/info/rfc5785>>.

- [RFC5925] Touch, J., Mankin, A., and R. Bonica, "The TCP Authentication Option", RFC 5925, DOI 10.17487/RFC5925, June 2010, <<https://www.rfc-editor.org/info/rfc5925>>.
- [RFC6125] Saint-Andre, P. and J. Hodges, "Representation and Verification of Domain-Based Application Service Identity within Internet Public Key Infrastructure Using X.509 (PKIX) Certificates in the Context of Transport Layer Security (TLS)", RFC 6125, DOI 10.17487/RFC6125, March 2011, <<https://www.rfc-editor.org/info/rfc6125>>.
- [RFC6234] Eastlake 3rd, D. and T. Hansen, "US Secure Hash Algorithms (SHA and SHA-based HMAC and HKDF)", RFC 6234, DOI 10.17487/RFC6234, May 2011, <<https://www.rfc-editor.org/info/rfc6234>>.
- [RFC6347] Rescorla, E. and N. Modadugu, "Datagram Transport Layer Security Version 1.2", RFC 6347, DOI 10.17487/RFC6347, January 2012, <<https://www.rfc-editor.org/info/rfc6347>>.
- [RFC7250] Wouters, P., Ed., Tschofenig, H., Ed., Gilmore, J., Weiler, S., and T. Kivinen, "Using Raw Public Keys in Transport Layer Security (TLS) and Datagram Transport Layer Security (DTLS)", RFC 7250, DOI 10.17487/RFC7250, June 2014, <<https://www.rfc-editor.org/info/rfc7250>>.
- [RFC7252] Shelby, Z., Hartke, K., and C. Bormann, "The Constrained Application Protocol (CoAP)", RFC 7252, DOI 10.17487/RFC7252, June 2014, <<https://www.rfc-editor.org/info/rfc7252>>.
- [RFC7525] Sheffer, Y., Holz, R., and P. Saint-Andre, "Recommendations for Secure Use of Transport Layer Security (TLS) and Datagram Transport Layer Security (DTLS)", BCP 195, RFC 7525, DOI 10.17487/RFC7525, May 2015, <<https://www.rfc-editor.org/info/rfc7525>>.
- [RFC7641] Hartke, K., "Observing Resources in the Constrained Application Protocol (CoAP)", RFC 7641, DOI 10.17487/RFC7641, September 2015, <<https://www.rfc-editor.org/info/rfc7641>>.
- [RFC7950] Bjorklund, M., Ed., "The YANG 1.1 Data Modeling Language", RFC 7950, DOI 10.17487/RFC7950, August 2016, <<https://www.rfc-editor.org/info/rfc7950>>.
- [RFC8126] Cotton, M., Leiba, B., and T. Narten, "Guidelines for Writing an IANA Considerations Section in RFCs", BCP 26, RFC 8126, DOI 10.17487/RFC8126, June 2017, <<https://www.rfc-editor.org/info/rfc8126>>.
- [RFC8132] van der Stok, P., Bormann, C., and A. Sehgal, "PATCH and FETCH Methods for the Constrained Application Protocol (CoAP)", RFC 8132, DOI 10.17487/RFC8132, April 2017, <<https://www.rfc-editor.org/info/rfc8132>>.

#### 14.2. Informative References

- [I-D.ietf-core-comi] Veillette, M., Stok, P., Pelov, A., and A. Bierman, "CoAP Management Interface", draft-ietf-core-comi-02 (work in progress), December 2017.
- [I-D.ietf-core-yang-cbor] Veillette, M., Pelov, A., Somaraju, A., Turner, R., and A. Minaburo, "CBOR Encoding of Data Modeled with YANG", draft-ietf-core-yang-cbor-05 (work in progress), August 2017.
- [I-D.ietf-dots-architecture] Mortensen, A., Andreasen, F., Reddy, T., christopher\_gray@cable.comcast.com, c., Compton, R., and N. Teague, "Distributed-Denial-of-Service Open Threat Signaling (DOTS) Architecture", draft-ietf-dots-architecture-05 (work in progress), October 2017.
- [I-D.ietf-dots-data-channel] Reddy, T., Boucadair, M., Nishizuka, K., Xia, L., Patil, P., Mortensen, A., and N. Teague, "Distributed Denial-of-Service Open Threat Signaling (DOTS) Data Channel", draft-ietf-dots-data-channel-11 (work in progress), December 2017.

[I-D.ietf-dots-requirements]  
Mortensen, A., Moskowitz, R., and T. Reddy, "Distributed Denial of Service (DDoS) Open Threat Signaling Requirements", ~~draft-ietf-dots-requirements-08~~ **draft-ietf-dots-requirements-09** (work in progress), December 2017.

[I-D.ietf-dots-use-cases]  
Dobbins, R., Migault, D., Fouant, S., Moskowitz, R., Teague, N., Xia, L., and K. Nishizuka, "Use cases for DDoS Open Threat Signaling", draft-ietf-dots-use-cases-09 (work in progress), November 2017.

[I-D.ietf-netmod-yang-tree-diagrams]  
Bjorklund, M. and L. Berger, "YANG Tree Diagrams", draft-~~ietf-netmod-yang-tree-diagrams-02~~ **ietf-netmod-yang-tree-diagrams-03** (work in progress), ~~October~~ **December** 2017.

[I-D.ietf-tls-dtls13]  
Rescorla, E., Tschofenig, H., and N. Modadugu, "The Datagram Transport Layer Security (DTLS) Protocol Version 1.3", draft-ietf-tls-dtls13-22 (work in progress), November 2017.

[I-D.ietf-tls-tls13]  
Rescorla, E., "The Transport Layer Security (TLS) Protocol Version 1.3", draft-ietf-tls-tls13-22 (work in progress), November 2017.

[proto\_numbers]  
"IANA, "Protocol Numbers", 2011,  
<<http://www.iana.org/assignments/protocol-numbers>>.

[RFC0791] Postel, J., "Internet Protocol", STD 5, RFC 791, DOI 10.17487/RFC0791, September 1981, <<https://www.rfc-editor.org/info/rfc791>>.

[RFC3022] Srisuresh, P. and K. Egevang, "Traditional IP Network Address Translator (Traditional NAT)", RFC 3022, DOI 10.17487/RFC3022, January 2001, <<https://www.rfc-editor.org/info/rfc3022>>.

[RFC3986] Berners-Lee, T., Fielding, R., and L. Masinter, "Uniform Resource Identifier (URI): Generic Syntax", STD 66, RFC 3986, DOI 10.17487/RFC3986, January 2005, <<https://www.rfc-editor.org/info/rfc3986>>.

**[RFC4086] Eastlake 3rd, D., Schiller, J., and S. Crocker, "Randomness Requirements for Security", BCP 106, RFC 4086, DOI 10.17487/RFC4086, June 2005, <<https://www.rfc-editor.org/info/rfc4086>>.**

[RFC4340] Kohler, E., Handley, M., and S. Floyd, "Datagram Congestion Control Protocol (DCCP)", RFC 4340, DOI 10.17487/RFC4340, March 2006, <<https://www.rfc-editor.org/info/rfc4340>>.

[RFC4632] Fuller, V. and T. Li, "Classless Inter-domain Routing (CIDR): The Internet Address Assignment and Aggregation Plan", BCP 122, RFC 4632, DOI 10.17487/RFC4632, August 2006, <<https://www.rfc-editor.org/info/rfc4632>>.

[RFC4732] Handley, M., Ed., Rescorla, E., Ed., and IAB, "Internet Denial-of-Service Considerations", RFC 4732, DOI 10.17487/RFC4732, December 2006, <<https://www.rfc-editor.org/info/rfc4732>>.

[RFC4787] Audet, F., Ed. and C. Jennings, "Network Address Translation (NAT) Behavioral Requirements for Unicast UDP", BCP 127, RFC 4787, DOI 10.17487/RFC4787, January 2007, <<https://www.rfc-editor.org/info/rfc4787>>.

**[RFC4941] Narten, T., Draves, R., and S. Krishnan, "Privacy Extensions for Stateless Address Autoconfiguration in IPv6", RFC 4941, DOI 10.17487/RFC4941, September 2007, <<https://www.rfc-editor.org/info/rfc4941>>.**

[RFC4960] Stewart, R., Ed., "Stream Control Transmission Protocol", RFC 4960, DOI 10.17487/RFC4960, September 2007, <<https://www.rfc-editor.org/info/rfc4960>>.

[RFC4987] Eddy, W., "TCP SYN Flooding Attacks and Common

- Mitigations", RFC 4987, DOI 10.17487/RFC4987, August 2007, <<https://www.rfc-editor.org/info/rfc4987>>.
- [RFC5077] Salowey, J., Zhou, H., Eronen, P., and H. Tschofenig, "Transport Layer Security (TLS) Session Resumption without Server-Side State", RFC 5077, DOI 10.17487/RFC5077, January 2008, <<https://www.rfc-editor.org/info/rfc5077>>.
- [RFC5389] Rosenberg, J., Mahy, R., Matthews, P., and D. Wing, "Session Traversal Utilities for NAT (STUN)", RFC 5389, DOI 10.17487/RFC5389, October 2008, <<https://www.rfc-editor.org/info/rfc5389>>.
- [RFC6146] Bagnulo, M., Matthews, P., and I. van Beijnum, "Stateful NAT64: Network Address and Protocol Translation from IPv6 Clients to IPv4 Servers", RFC 6146, DOI 10.17487/RFC6146, April 2011, <<https://www.rfc-editor.org/info/rfc6146>>.
- [RFC6296] Wasserman, M. and F. Baker, "IPv6-to-IPv6 Network Prefix Translation", RFC 6296, DOI 10.17487/RFC6296, June 2011, <<https://www.rfc-editor.org/info/rfc6296>>.
- [RFC6555] Wing, D. and A. Yourtchenko, "Happy Eyeballs: Success with Dual-Stack Hosts", RFC 6555, DOI 10.17487/RFC6555, April 2012, <<https://www.rfc-editor.org/info/rfc6555>>.
- [RFC6724] Thaler, D., Ed., Draves, R., Matsumoto, A., and T. Chown, "Default Address Selection for Internet Protocol Version 6 (IPv6)", RFC 6724, DOI 10.17487/RFC6724, September 2012, <<https://www.rfc-editor.org/info/rfc6724>>.
- [RFC6887] Wing, D., Ed., Cheshire, S., Boucadair, M., Penno, R., and P. Selkirk, "Port Control Protocol (PCP)", RFC 6887, DOI 10.17487/RFC6887, April 2013, <<https://www.rfc-editor.org/info/rfc6887>>.
- [RFC6888] Perreault, S., Ed., Yamagata, I., Miyakawa, S., Nakagawa, A., and H. Ashida, "Common Requirements for Carrier-Grade NATs (CGNs)", BCP 127, RFC 6888, DOI 10.17487/RFC6888, April 2013, <<https://www.rfc-editor.org/info/rfc6888>>.
- [RFC7030] Pritikin, M., Ed., Yee, P., Ed., and D. Harkins, Ed., "Enrollment over Secure Transport", RFC 7030, DOI 10.17487/RFC7030, October 2013, <<https://www.rfc-editor.org/info/rfc7030>>.
- [RFC7049] Bormann, C. and P. Hoffman, "Concise Binary Object Representation (CBOR)", RFC 7049, DOI 10.17487/RFC7049, October 2013, <<https://www.rfc-editor.org/info/rfc7049>>.
- [RFC7413] Cheng, Y., Chu, J., Radhakrishnan, S., and A. Jain, "TCP Fast Open", RFC 7413, DOI 10.17487/RFC7413, December 2014, <<https://www.rfc-editor.org/info/rfc7413>>.
- [RFC7452] Tschofenig, H., Arkko, J., Thaler, D., and D. McPherson, "Architectural Considerations in Smart Object Networking", RFC 7452, DOI 10.17487/RFC7452, March 2015, <<https://www.rfc-editor.org/info/rfc7452>>.
- [RFC7589] Badra, M., Luchuk, A., and J. Schoenwaelder, "Using the NETCONF Protocol over Transport Layer Security (TLS) with Mutual X.509 Authentication", RFC 7589, DOI 10.17487/RFC7589, June 2015, <<https://www.rfc-editor.org/info/rfc7589>>.
- [RFC7918] Langley, A., Modadugu, N., and B. Moeller, "Transport Layer Security (TLS) False Start", RFC 7918, DOI 10.17487/RFC7918, August 2016, <<https://www.rfc-editor.org/info/rfc7918>>.
- [RFC7924] Santesson, S. and H. Tschofenig, "Transport Layer Security (TLS) Cached Information Extension", RFC 7924, DOI 10.17487/RFC7924, July 2016, <<https://www.rfc-editor.org/info/rfc7924>>.
- [RFC7942] Sheffer, Y. and A. Farrel, "Improving Awareness of Running Code: The Implementation Status Section", BCP 205, RFC 7942, DOI 10.17487/RFC7942, July 2016, <<https://www.rfc-editor.org/info/rfc7942>>.
- [RFC7951] Lhotka, L., "JSON Encoding of Data Modeled with YANG", RFC 7951, DOI 10.17487/RFC7951, August 2016, <<https://www.rfc-editor.org/info/rfc7951>>.



[RFC8085] Eggert, L., Fairhurst, G., and G. Shepherd, "UDP Usage Guidelines", BCP 145, RFC 8085, DOI 10.17487/RFC8085, March 2017, <<https://www.rfc-editor.org/info/rfc8085>>.

#### Authors' Addresses

Tirumaleswar Reddy  
McAfee, Inc.  
Embassy Golf Link Business Park  
Bangalore, Karnataka 560071  
India

Email: [kondtir@gmail.com](mailto:kondtir@gmail.com)

Mohamed Boucadair  
Orange  
Rennes 35000  
France

Email: [mohamed.boucadair@orange.com](mailto:mohamed.boucadair@orange.com)  
Prashanth Patil  
Cisco Systems, Inc.

Email: [praspati@cisco.com](mailto:praspati@cisco.com)

Andrew Mortensen  
Arbor Networks, Inc.  
2727 S. State St  
Ann Arbor, MI 48104  
United States

Email: [amortensen@arbor.net](mailto:amortensen@arbor.net)

Nik Teague  
Verisign, Inc.  
United States

Email: [nteague@verisign.com](mailto:nteague@verisign.com)